

Kubernetes resources sharing




DMITRII DOLGOV

22-10-2018



- CPU
- CPU cache
- Storage IO
- Network
- Memory
- Hugetlb

container = cgroup + namespace




CPU

directly manageable

requests → `cpu.share`

limits → `cpu.cfs_period` & `cpu.cfs_quota_us`


Share



Share



Bandwidth



Bandwidth



Bandwidth accounting

```
# from /proc/sys/kernel/  
sched_cfs_bandwidth_slice_us  
# default=5ms
```

Allocatable

Capacity:

cpu:	16
hugepages-1Gi:	0
hugepages-2Mi:	0
memory:	65947396Ki

Allocatable:

cpu:	15800m
hugepages-1Gi:	0
hugepages-2Mi:	0
memory:	65388292Ki

Exclusive CPU

cpu management policy

kube-reserved

guaranteed

integer quantity

cpuset.cpus

cpuset.cpuset.cpu_exclusive

cpuset.mems ?

Cache

```
# COS1 4 cache ways, COS 2 next 8 cache ways  
=> pqos -e "llc:1=0x000f;llc:2=0x0ff0;"  
SOCKET 0 L3CA COS1 => MASK 0xf  
SOCKET 0 L3CA COS2 => MASK 0xff0  
Allocation configuration altered.
```

Cache

=> pqos -s

L3CA COS definitions **for** Socket **0**:

L3CA COS0 => MASK 0xffff

L3CA COS1 => MASK 0xf

L3CA COS2 => MASK 0xff0

L3CA COS3 => MASK 0xfff

Core information **for** socket **0**:

Core **0**, L2ID **0**, L3ID **0** => COS0

Core **1**, L2ID **1**, L3ID **0** => COS0

Core **2**, L2ID **0**, L3ID **0** => COS0

Core **3**, L2ID **1**, L3ID **0** => COS0

Memory

directly manageable

requests ↗ `memory.soft_limit_in_bytes`

limits → `memory.limit_in_bytes` (OOM)

`memory.kmem.limit_in_bytes`

best efforts (not everything is accounted)

Memory reclaim

only under the memory pressure

```
root@k8s-node-2:/home/vagrant# ./page_reclaim.py
```

Attaching...

Listening...

Detaching...

[7382] postgres: 928K

[7138] postgres: 152K

[7136] postgres: 180K

[7468] postgres: 72M

[7464] postgres: 57M

[5451] postgres: 1M

Writeback (cgroup v1)

```
/* vmscan.c */  
/* The normal page dirty throttling mechanism  
 * in balance_dirty_pages() is completely broken  
 * with the legacy memcg and direct stalling in  
 * shrink_page_list() is used for throttling instead,  
 * which lacks all the niceties such as fairness,  
 * adaptive pausing, bandwidth proportional  
 * allocation and configurability.  
 */  
static bool sane_reclaim(struct scan_control *sc)
```

Writeback monitoring

```
=> perf record -e writeback:writeback_written  
kworker/u8:1 5816.288044: nr_pages=101429  
kworker/u8:1 5816.288129: nr_pages=9223372036854775789  
kworker/u8:3 5817.312319: nr_pages=101457
```

Writeback monitoring

```
# pgbench insert
=> ./io_timeouts.py -p bin/postgres
Attaching...
Listening...
Detaching...
[18335] END: MAX_SCHEDULE_TIMEOUT
[18333] END: MAX_SCHEDULE_TIMEOUT
[18331] END: MAX_SCHEDULE_TIMEOUT
[18318] truncate pgbench_history: MAX_SCHEDULE_TIMEOUT
```

Huge pages

directly manageable

transparent vs classic

isolation only per pod

no soft limits or reclaim (SIGBUS)

TLB misses are faster and less frequent

memory leaks (but PG is good)

Huge pages

huge_pages on

Samples: 832K of event 'dTLB-load-misses'

Event count (approx.): **640614445** : ~19% less

Samples: 736K of event 'dTLB-store-misses'

Event count (approx.): **72447300** : ~29% less

huge_pages off

Samples: 894K of event 'dTLB-load-misses'

Event count (approx.): **784439650**

Samples: 822K of event 'dTLB-store-misses'

Event count (approx.): **101471557**

Storage IO

blkio.weight

blkio.throttle.*

Not used in K8S

sane_behavior

cpuset.cpus

cpuset.cpuset.cpu_exclusive

cpuset.mems ?

IO scheduler

```
=> cat /sys/block/xvdcj/queue/scheduler  
[mq-deadline] kyber bfq none
```

IO scheduler

BFQ distributes the bandwidth of the device among all processes according to their weights, regardless of the device parameters and with any workload.

IO scheduler

The Kyber I/O scheduler is a low-overhead scheduler suitable for multiqueue and other fast devices. Given target latencies for reads and synchronous writes, it will self-tune queue depths to achieve that goal.

Network

not directly
network class
traffic control

Kernel noise

Futex

Compaction

Readahead


Filesystem


...

Questions?

 github.com/erthalion

 [@erthalion](https://twitter.com/erthalion)

 dmitrii.dolgov at zalando dot de

 9erthalion6 at gmail dot com