

K8S RESOURCE SHARING

FOR POSTGRESQL
EXPERTS

DMITRY DOLGOV

03-09-2019





patroni & postgres-operator

```
resources:
  requests:
    cpu: "1000m"
    memory: "1GB"
  limits:
    cpu: "2000m"
    memory: "2GB"
```



Allocatable

Capacity:

cpu:	16
hugepages-1Gi:	0
hugepages-2Mi:	0
memory:	65947396Ki

Allocatable:

cpu:	15800m
hugepages-1Gi:	0
hugepages-2Mi:	0
memory:	65388292Ki

Allocatable

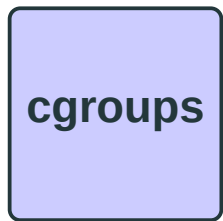
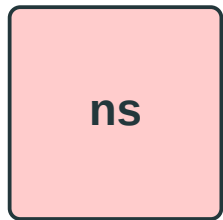
Capacity:


cpu:	16
hugepages-1Gi:	0
hugepages-2Mi:	0
memory:	65947396Ki
Allocatable:	
cpu:	15800m
hugepages-1Gi:	0
hugepages-2Mi:	0
memory:	65388292Ki

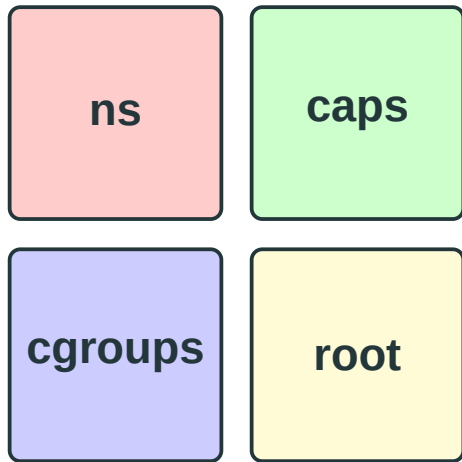
K8S appliance?




cgroups









→ docker

→ rkt

→ LXC


→ ...

→ podman

→ cgroup v1

→ cgroup v2

→ ...



resources:

requests:

cpu: "1000m"

memory: "1GB"

limits:

cpu: "2000m"


memory: "2GB"



cpu share

cfs quota

Share



Share



Quota



Quota



Bandwidth accounting

```
# from /proc/sys/kernel/  
sched_cfs_bandwidth_slice_us  
# default=5ms
```

Exclusive CPU

- cpu management policy
- kube-reserved
- guaranteed (integer quantity)
- cpuset.cpus
- cpuset.cpuset.cpu_exclusive

resources:

requests:

cpu: "1000m"

memory: "1GB"

limits:

cpu: "2000m"

memory: "2GB"



soft mem limit

hard mem limit

resources:

requests:

cpu: "1000m"

memory: "1GB"

limits:

cpu: "2000m"

memory: "2GB"

~~soft mem limit~~

hard mem limit



Memory reclaim

only under the memory pressure

=> `page_reclaim.py --container 89c33bb3133f`

[7382] postgres: 928K

[7138] postgres: 152K

[7136] postgres: 180K

[7468] postgres: 72M

[7464] postgres: 57M

[5451] postgres: 1M

Memory

- best efforts
- requests for QoS and oom_adj
- memory.kmem.limit_in_bytes
- MEMCG_CHARGE_BATCH (32 pages)
- shared memory CoW

resources:

requests:

cpu: "1000m"

memory: "1GB"

hugepages-1Gi: "1GB"

hugepages-2Mi: "1GB"

limits:

cpu: "2000m"

memory: "2GB"

hugepages-1Gi: "2GB"

hugepages-2Mi: "2GB"

resources:

requests:

cpu: "1000m"

memory: "1GB"

hugepages-1Gi: "1GB"

hugepages-2Mi: "1GB"

limits:

cpu: "2000m"

memory: "2GB"

hugepages-1Gi: "2GB"

hugepages-2Mi: "2GB"

→ classic

resources:

requests:

cpu: "1000m"

memory: "1GB"

hugepages-1Gi: "1GB"

hugepages-2Mi: "1GB"

limits:

cpu: "2000m"

memory: "2GB"

hugepages-1Gi: "2GB"

hugepages-2Mi: "2GB"

→ classic

→ transparent

resources:

requests:

cpu: "1000m"

memory: "1GB"

hugepages-1Gi: "1GB"

hugepages-2Mi: "1GB"

limits:

cpu: "2000m"

memory: "2GB"

hugepages-1Gi: "2GB"

hugepages-2Mi: "2GB"

→ classic

→ transparent

→ isolation only per pod

resources:

requests:

cpu: "1000m"

memory: "1GB"

hugepages-1Gi: "1GB"

hugepages-2Mi: "1GB"

limits:

cpu: "2000m"

memory: "2GB"

hugepages-1Gi: "2GB"

hugepages-2Mi: "2GB"

→ classic

→ transparent

→ isolation only per pod

→ no soft limits or
reclaim (SIGBUS)

Huge pages

```
# perf record -e dTLB-loads,dTLB-stores -p PID
```

```
# huge_pages on
```

```
Samples: 832K of event 'dTLB-load-misses'
```

```
Event count (approx.): 640614445 : ~19% less
```

```
Samples: 736K of event 'dTLB-store-misses'
```

```
Event count (approx.): 72447300 : ~29% less
```

```
# huge_pages off
```

```
Samples: 894K of event 'dTLB-load-misses'
```

```
Event count (approx.): 784439650
```

```
Samples: 822K of event 'dTLB-store-misses'
```

```
Event count (approx.): 101471557
```

Huge pages

```
# perf record -e dTLB-loads,dTLB-stores -p PID
```

```
# huge_pages on
```

```
Samples: 832K of event 'dTLB-load-misses'
```

```
Event count (approx.): 640614445 : ~19% less
```

```
Samples: 736K of event 'dTLB-store-misses'
```

```
Event count (approx.): 72447300 : ~29% less
```

```
# huge_pages off
```

```
Samples: 894K of event 'dTLB-load-misses'
```

```
Event count (approx.): 784439650
```

```
Samples: 822K of event 'dTLB-store-misses'
```

```
Event count (approx.): 101471557
```

LLCache

=> pqos -s

L3CA COS definitions **for** Socket **0**:

L3CA COS0 => MASK 0xffff

L3CA COS1 => MASK 0xf

L3CA COS2 => MASK 0xff0

L3CA COS3 => MASK 0xfff

Core information **for** socket **0**:

Core **0**, L2ID **0**, L3ID 0 => COS0

Core **1**, L2ID **1**, L3ID 0 => COS0

Core **2**, L2ID **0**, L3ID 0 => COS0

Core **3**, L2ID **1**, L3ID 0 => COS0

LLCache

```
# COS1 4 cache ways, COS 2 next 8 cache ways  
=> pqos -e "llc:1=0x000f;llc:2=0x0ff0;"  
SOCKET 0 L3CA COS1 => MASK 0xf  
SOCKET 0 L3CA COS2 => MASK 0xff0  
Allocation configuration altered.
```

Block I/O

The creation of the io.latency block I/O controller

Block I/O

→ sharing a disk between users in Linux is awful

The creation of the io.latency block I/O controller

Block I/O

- sharing a disk between users in Linux is awful
- throttling puts pressure somewhere else

The creation of the io.latency block I/O controller

Block I/O

- sharing a disk between users in Linux is awful
- throttling puts pressure somewhere else
- lack of trivially observable cost metric

The creation of the io.latency block I/O controller

Writeback monitoring

```
=> perf record -e writeback:writeback_written
```

```
kworker/u8:1 reason=periodic    nr_pages=101429  
kworker/u8:1 reason=background nr_pages=MAX_ULONG  
kworker/u8:3 reason=periodic    nr_pages=101457
```

Writeback monitoring

```
# pgbench insert workload
```

```
=> io_timeouts.py bin/postgres
```

```
[18335] END: MAX_SCHEDULE_TIMEOUT
```

```
[18333] END: MAX_SCHEDULE_TIMEOUT
```

```
[18331] END: MAX_SCHEDULE_TIMEOUT
```

```
[18318] truncate pgbench_history: MAX_SCHEDULE_TIMEOUT
```

blkio cgroup v1


- Direct IO oriented
- no relationships between controllers
- writeback = memory + blkio controllers
- writeback accounted to the root cgroup

<https://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux.git>


IO scheduler

```
=> cat /sys/block/xvdcj/queue/scheduler  
[mq-deadline] kyber bfq none
```

I/O sched



blkmq



```
=> blk_mq.py --container 89c33bb3133f
```



latency (us)	: count	distribution
16 -> 31	: 0	
32 -> 63	: 19	***
64 -> 127	: 27	****
128 -> 255	: 6	*
256 -> 511	: 8	*
512 -> 1023	: 17	***
1024 -> 2047	: 40	*****
2048 -> 4095	: 126	*****
4096 -> 8191	: 144	*****
8192 -> 16383	: 222	*****
16384 -> 32767	: 120	*****
32768 -> 65535	: 44	*****



index : kernel/git/next/linux-next.git

The linux-next integration testing tree

[about](#) [summary](#) [refs](#) [log](#) [tree](#) **[commit](#)** [diff](#) [stats](#)

author  Tejun Heo <tj@kernel.org> 2019-08-28 15:05:58 -0700
committer  Jens Axboe <axboe@kernel.dk> 2019-08-28 21:17:12 -0600
commit [7caa47151ab2e644dd221f741ec7578d9532c9a3](#) (patch)
tree [f5ffe39d84924c03fb72f927ab420e3ca6a629ec](#)
parent [6f816b4b746c2241540e537682d30d8e9997d674](#) (diff)
download [linux-next-7caa47151ab2e644dd221f741ec7578d9532c9a3.tar.gz](#)

blkcg: implement blk-iocost

This patchset implements IO cost model based work-conserving proportional controller.

Network

- vlanx
- network class
- traffic control


Noise


- Futex
- Compaction
- Readahead
- Filesystem
- ...

Questions?

 github.com/erthalion

 [@erthalion](https://twitter.com/erthalion)

 [dmitrii.dolgov at zalando dot de](mailto:dmitrii.dolgov@zalando.de)

 [9erthalion6 at gmail dot com](mailto:9erthalion6@gmail.com)