

POSTGRESQL AT LOW LEVEL

STAY CURIOUS!

DMITRY DOLGOV

20-09-2019





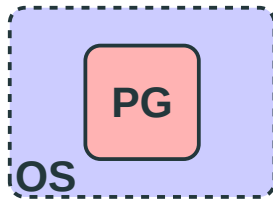
patroni & postgres-operator


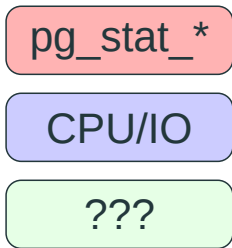
pg_stat_*

PG

pg_stat_*

CPU/IO



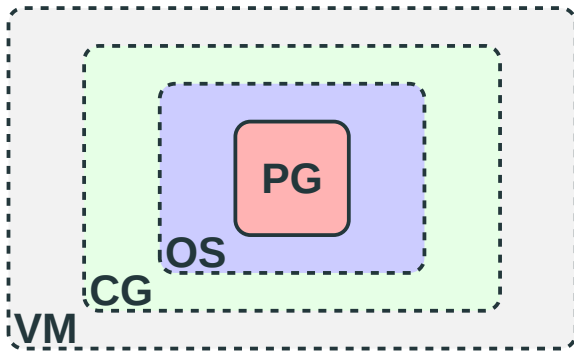


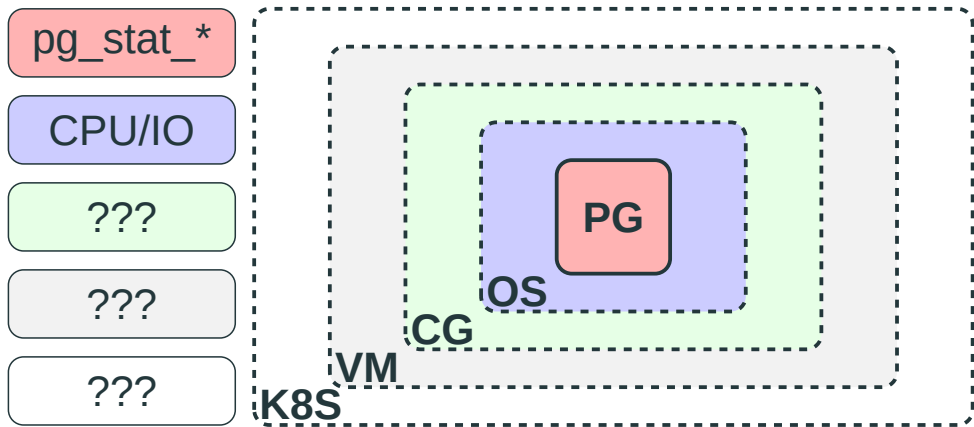
pg_stat_*

CPU/IO

???

???







All right, then. Keep your secrets

Plan?

A bit chaotic



Info sources

source code

strace/GDB/Perf

procfs/sysfs

BPF/eBPF/BCC

Shared memory

```
ERROR:  could not resize shared memory segment  
"/PostgreSQL.699663942" to 50438144 bytes:  
No space left on device
```

```
# strace -k -p PID
openat(AT_FDCWD, "/dev/shm/PostgreSQL.62223175"
ftruncate(176, 50438144) = 0
fallocate(176, 0, 0, 50438144) = -1 ENOSPC
> libc-2.27.so(posix_fallocate+0x16) [0x114f76]
> postgres(dsm_create+0x67) [0x377067]
...
> postgres(ExecInitParallelPlan+0x360) [0x254a80]
> postgres(ExecGather+0x495) [0x269115]
> postgres(standard_ExecutorRun+0xfd) [0x25099d]
...
> postgres(exec_simple_query+0x19f) [0x39afdf]
```

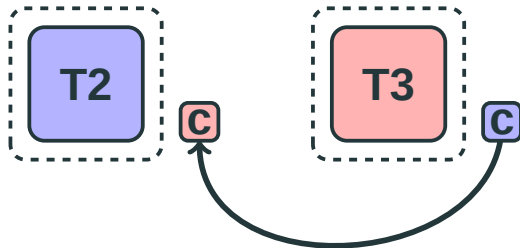
```
# strace -k -p PID
openat(AT_FDCWD, "/dev/shm/PostgreSQL.62223175"
ftruncate(176, 50438144)                = 0
fallocate(176, 0, 0, 50438144)          = -1 ENOSPC
> libc-2.27.so(posix_fallocate+0x16) [0x114f76]
> postgres(dsm_create+0x67) [0x377067]
...
> postgres(ExecInitParallelPlan+0x360) [0x254a80]
> postgres(ExecGather+0x495) [0x269115]
> postgres(standard_ExecutorRun+0xfd) [0x25099d]
...
> postgres(exec_simple_query+0x19f) [0x39afdf]
```

vDSO

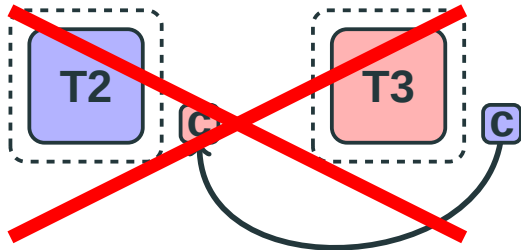
```
# strace -k -p PID on XEN  
gettimeofday({tv_sec=1550586520, tv_usec=313499}, NULL) = 0  
> [vdso]() [0xef0]
```

Two frequently used system calls are 77% slower on AWS EC2

Scheduling



Scheduling





Andres Freund: New intel MDS vulnerability mitigations cause measurable slowdown

MDS

# Children	Self	Symbol
#
71.06%	0.00%	[.] __libc_start_main
71.06%	0.00%	[.] PostmasterMain
56.82%	0.14%	[.] exec_simple_query
25.19%	0.06%	[k] entry_SYSCALL_64_after_hwframe
25.14%	0.29%	[k] do_syscall_64
23.60%	0.14%	[.] standard_ExecutorRun

MDS

# Children	Self	Symbol
#
71.06%	0.00%	[.] __libc_start_main
71.06%	0.00%	[.] PostmasterMain
56.82%	0.14%	[.] exec_simple_query
25.19%	0.06%	[k] entry_SYSCALL_64_after_hwframe
25.14%	0.29%	[k] do_syscall_64
23.60%	0.14%	[.] standard_ExecutorRun

MDS

```
# Percent      Disassembly of kcore for cycles
# .....
  0.01% :      nopl      0x0(%rax,%rax,1)
28.94% :      verw      0xffe9e1(%rip)
  0.55% :      pop      %rbx
  3.24% :      pop      %rbp
```

MDS

```
# Percent      Disassembly of kcore for cycles
# .....
  0.01% :      nopl      0x0(%rax,%rax,1)
 28.94% :      verw     0xffe9e1(%rip)
  0.55% :      pop      %rbx
  3.24% :      pop      %rbp
```

MDS

#	Overhead	Symbol
#
	25.19%	[k] native_safe_halt

MDS

```
static inline __cpuidle void native_safe_halt(void)
{
    mds_idle_clear_cpu_buffers();
    asm volatile("sti; hlt" : : : "memory");
}
```


MDS

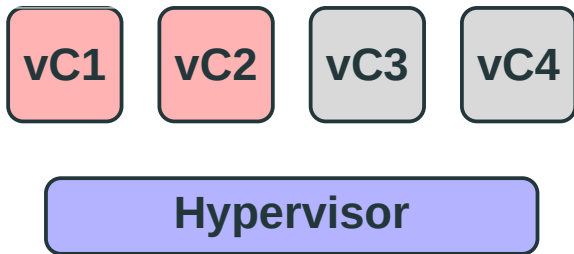
```
static inline __cpuidle void native_safe_halt(void)
{
    mds_idle_clear_cpu_buffers();
    asm volatile("sti; hlt" : : : "memory");
}
```

VM


- Lock holder preemption problem
- Lock waiter preemption problem
- Intel PLE (pause loop exiting)
- PLE_Gap, PLE_Window

Intel® 64 and IA-32 Architectures Software Developer's Manual, Vol. 3


vCPU



vCPU



vCPU




```
# latency average = 17.782 ms  
=> modprobe kvm-intel ple_gap=128  
=> perf record -e kvm:kvm_exit  
reason PAUSE_INSTRUCTION 306795
```


```
# latency average = 17.782 ms  
=> modprobe kvm-intel ple_gap=128  
=> perf record -e kvm:kvm_exit  
reason PAUSE_INSTRUCTION 306795
```


```
# latency average = 16.858 ms  
=> modprobe kvm-intel ple_gap=0  
=> perf record -e kvm:kvm_exit  
reason PAUSE_INSTRUCTION 0
```


```
# latency average = 17.782 ms  
=> modprobe kvm-intel ple_gap=128  
=> perf record -e kvm:kvm_exit  
reason PAUSE_INSTRUCTION 306795
```


```
# latency average = 16.858 ms  
=> modprobe kvm-intel ple_gap=0  
=> perf record -e kvm:kvm_exit  
reason PAUSE_INSTRUCTION 0
```









Huge pages

transparent vs classic

TLB misses are faster and less frequent

Huge pages

```
# perf record -e dTLB-loads,dTLB-stores -p PID
```

```
# huge_pages on
```

```
Samples: 832K of event 'dTLB-load-misses'
```

```
Event count (approx.): 640614445 : ~19% less
```

```
Samples: 736K of event 'dTLB-store-misses'
```

```
Event count (approx.): 72447300 : ~29% less
```

```
# huge_pages off
```

```
Samples: 894K of event 'dTLB-load-misses'
```

```
Event count (approx.): 784439650
```

```
Samples: 822K of event 'dTLB-store-misses'
```

```
Event count (approx.): 101471557
```

Huge pages

```
# perf record -e dTLB-loads,dTLB-stores -p PID
```

```
# huge_pages on
```

```
Samples: 832K of event 'dTLB-load-misses'
```

```
Event count (approx.): 640614445 : ~19% less
```

```
Samples: 736K of event 'dTLB-store-misses'
```

```
Event count (approx.): 72447300 : ~29% less
```


```
# huge_pages off
```

```
Samples: 894K of event 'dTLB-load-misses'
```

```
Event count (approx.): 784439650
```

```
Samples: 822K of event 'dTLB-store-misses'
```

```
Event count (approx.): 101471557
```





**And now for something
completely different**


Userspace

Bytecode

Userspace



Userspace



github.com/iovisor/bcc/

github.com/erthalion/postgres-bcc

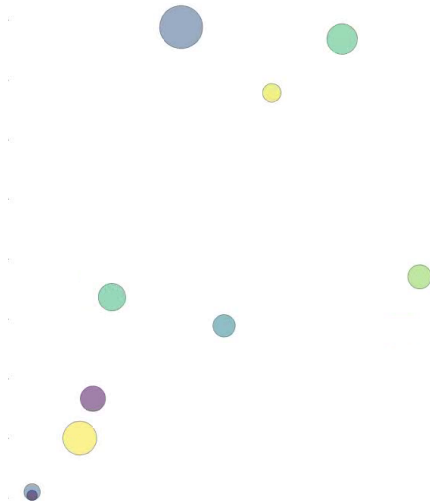
Cache

=> llcache_per_query.py bin/postgres

PID	QUERY	CPU	REFERENCE	MISS	HIT%
9720	UPDATE pgbench_tellers ...	0	2000	1000	50.00%
9720	SELECT abalance FROM ...	2	2000	100	95.00%
...					

Total References: **3303100** Total Misses: **599100** Hit Rate: **81.86%**

Shared buffers access



Writeback

```
=> perf record -e writeback:writeback_written
```

```
kworker/u8:1 reason=periodic    nr_pages=101429  
kworker/u8:1 reason=background  nr_pages=MAX_ULONG  
kworker/u8:3 reason=periodic    nr_pages=101457
```


Writeback

```
# pgbench insert workload
```

```
=> io_timeouts.py bin/postgres
```

```
[18335] END: MAX_SCHEDULE_TIMEOUT
```

```
[18333] END: MAX_SCHEDULE_TIMEOUT
```

```
[18331] END: MAX_SCHEDULE_TIMEOUT
```

```
[18318] truncate pgbench_history: MAX_SCHEDULE_TIMEOUT
```

Kubernetes

resources:

requests:

memory: "64Mi"

cpu: "250m"

limits:

memory: "128Mi"

cpu: "500m"

Kubernetes

resources:

requests:

memory: "64Mi"

cpu: "250m"

limits:


memory: "128Mi"

cpu: "500m"

soft_limits_in_bytes

limits_in_bytes

Evil ha-ha-ha



Kubernetes

resources:

requests:

memory: "64Mi"

cpu: "250m"

limits:

memory: "128Mi"

cpu: "500m"

~~soft_limits_in_bytes~~



limits_in_bytes

Memory reclaim

only under the memory pressure

=> `page_reclaim.py --container 89c33bb3133f`

[7382] postgres: 928K

[7138] postgres: 152K

[7136] postgres: 180K

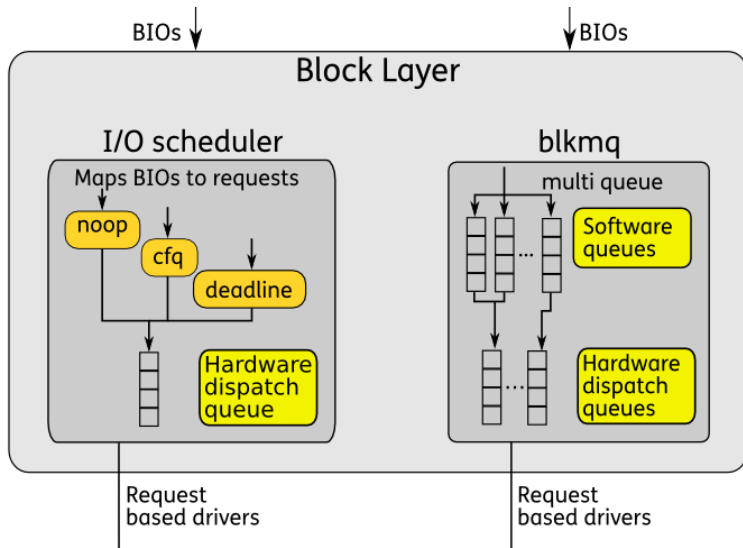
[7468] postgres: 72M

[7464] postgres: 57M

[5451] postgres: 1M

IO scheduler

```
=> cat /sys/block/xvdcj/queue/scheduler  
[mq-deadline] kyber bfq none
```



Werner Fischer


```
=> trace 'sbitmap_queue_resize "%d", arg2'
```

PID	TID	COMM	FUNC	-
22581	22581	postgres	sbitmap_queue_resize	74
22581	22581	postgres	sbitmap_queue_resize	51
22581	22581	postgres	sbitmap_queue_resize	80
22581	22581	postgres	sbitmap_queue_resize	45
22581	22581	postgres	sbitmap_queue_resize	81
22581	22581	postgres	sbitmap_queue_resize	47

```
=> blk_mq.py --container 89c33bb3133f
```



latency (us)	:	count	distribution
16 -> 31	:	0	
32 -> 63	:	19	***
64 -> 127	:	27	****
128 -> 255	:	6	*
256 -> 511	:	8	*
512 -> 1023	:	17	***
1024 -> 2047	:	40	*****
2048 -> 4095	:	126	*****
4096 -> 8191	:	144	*****
8192 -> 16383	:	222	*****
16384 -> 32767	:	120	*****
32768 -> 65535	:	44	*****



index : kernel/git/next/linux-next.git

The linux-next integration testing tree

[about](#) [summary](#) [refs](#) [log](#) [tree](#) **[commit](#)** [diff](#) [stats](#)

author  Tejun Heo <tj@kernel.org> 2019-08-28 15:05:58 -0700
committer  Jens Axboe <axboe@kernel.dk> 2019-08-28 21:17:12 -0600
commit [7caa47151ab2e644dd221f741ec7578d9532c9a3](#) (patch)
tree [f5ffe39d84924c03fb72f927ab420e3ca6a629ec](#)
parent [6f816b4b746c2241540e537682d30d8e9997d674](#) (diff)
download [linux-next-7caa47151ab2e644dd221f741ec7578d9532c9a3.tar.gz](#)

blkcg: implement blk-iocost

This patchset implements IO cost model based work-conserving proportional controller.

How to run?

```
# bcc + postgres-bcc
```

```
CONFIG_BPF=y
```

```
CONFIG_BPF_SYSCALL=y
```

```
CONFIG_NET_CLS_BPF=m
```

```
CONFIG_NET_ACT_BPF=m
```

```
CONFIG_BPF_JIT=y
```

```
CONFIG_BPF_EVENTS=y
```

```
debugfs on /sys/kernel/debug type debugfs (rw)
```

How to run: container?

```
# sometimes you also need to let perf know  
# where to find debugging symbols, e.g. copy  
# from /usr/lib/.debug/  
docker run  
    --privileged  
    --net=container:<container-id>  
    --ipc=container:<container-id>
```

How to run: K8S?

spec:

serviceAccountName: "bcc"

hostPID: true

containers:

- **name:** "bcc"

securityContext:

privileged: true

4 * 65536 + 14 * 256 + 96

=> export BCC_LINUX_VERSION_CODE **265824**

How to break?

```
# unsafe access
```

```
=> perf probe -x bin/postgres --funcs
```

```
=> perf probe -x bin/postgres 'ExecCallTriggerFunc trigdata->?'
```

```
=> perf record probe_postgres:ExecCallTriggerFunc
```

How to break?

non interruptible sleep

=> perf probe -x bin/postgres --funcs

=> perf probe -x bin/postgres 'XLogInsertRecord fpw_lsn'

How to break?

iovisor / bcc

Watch 410 ★ Uns

<> Code ⓘ Issues 266 🔗 Pull requests 21 📁 Projects 0 📖 Wiki 📊 Insights

Ubuntu xenial kernel panic in bpf_map_update_elem using ext4slower #1678


Closed stefreak opened this issue on Apr 12, 2018 · 13 comments


Questions?

 github.com/erthalion

 github.com/erthalion/postgres-bcc

 [@erthalion](https://twitter.com/erthalion)

 [dmitrii.dolgov at zalando dot de](mailto:dmitrii.dolgov@zalando.de)

 [9erthalion6 at gmail dot com](mailto:9erthalion6@gmail.com)