zalando

NOSQL FOR POSTGRESQL

BEST PRACTICES

DMITRY DOLGOV

10-25-2017

zalando

zalando

➔ Jsonb internals and performance-related factors

zalando

- ➔ Jsonb internals and performance-related factors
- ➔ Tricky queries

zalando

- ➔ Jsonb internals and performance-related factors
- ➔ Tricky queries
- ➔ Benchmarks

- ➔ Jsonb internals and performance-related factors
- ➔ Tricky queries
- ➔ Benchmarks
- ➔ How to shoot yourself in the foot

zalando

Internals

Performance-related factors

zalando

Performance-related factors

➔ On-disk representation

zalando

Performance-related factors

- ➔ On-disk representation
- ➔ In-memory representation

zalando

Performance-related factors

- ➔ On-disk representation
- ➔ In-memory representation
- ➔ Indexing support

zalando

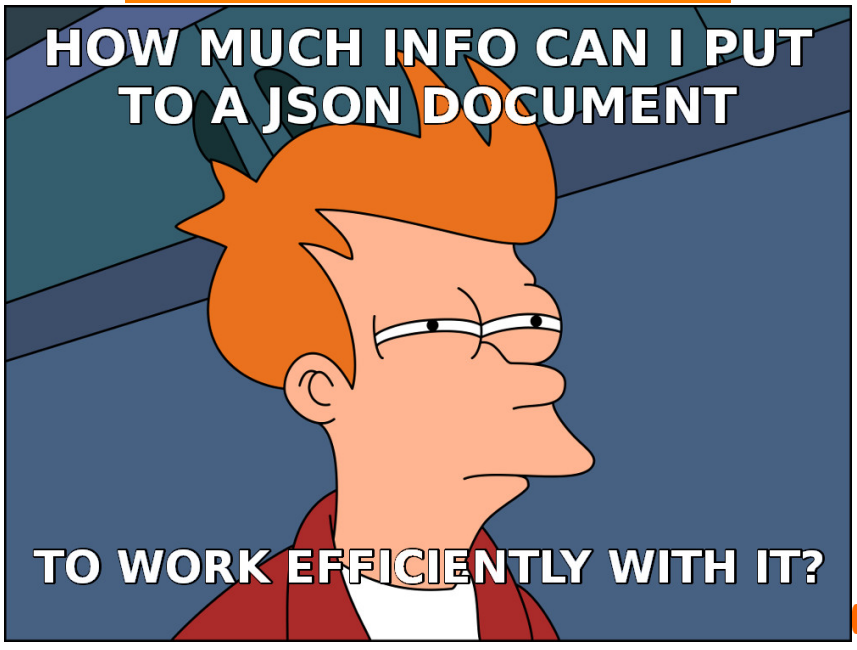
HOW MUCH INFO CAN I PUT TO A JSON DOCUMENT
TO WORK EFFICIENTLY WITH IT?

zalando

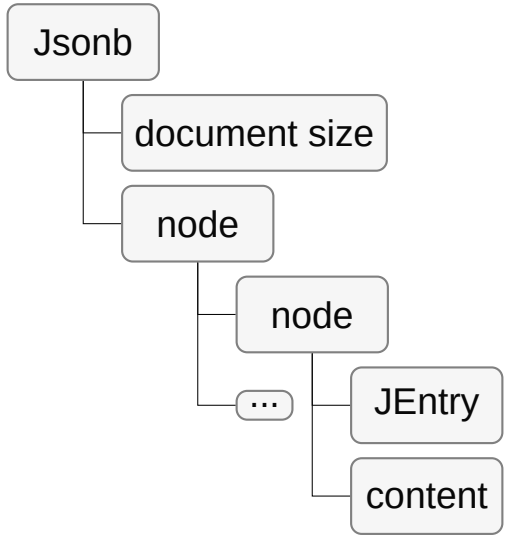

zalando

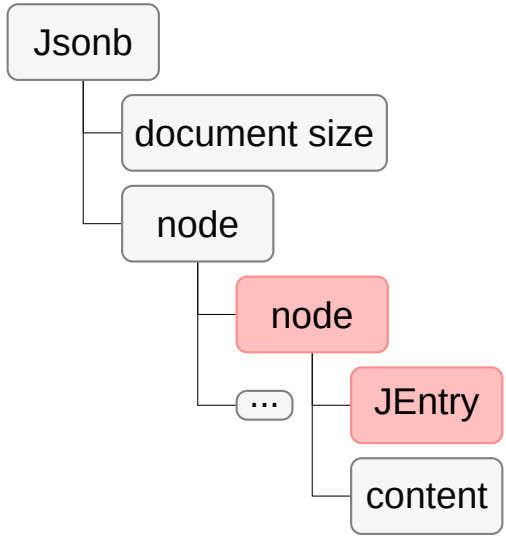

zalando

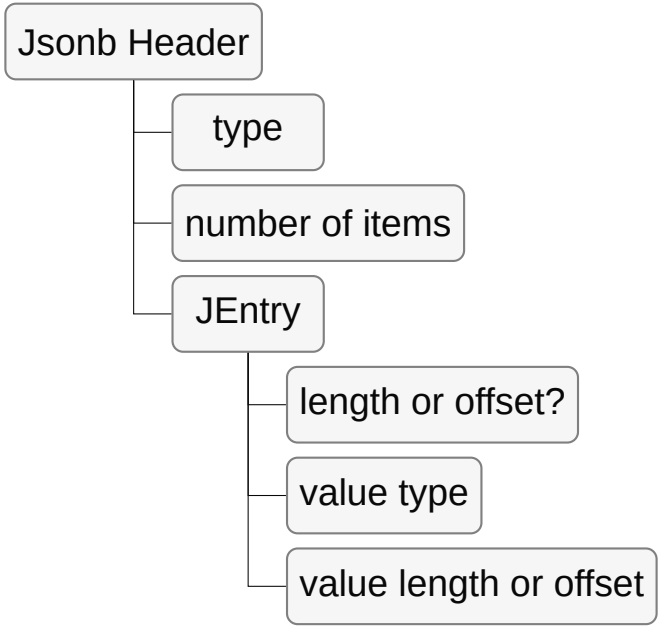
Jsonb Header
type
number of items
JEntry
length or offset?
value type
value length or offset

zalando

JB_OFFSET_STRIDE

- JEntry may contains a value lenght or offset
- Offset = access speed
- Length = compressibility
- Every **JB_OFFSET_STRIDE**’th JEntry contains an offset
- Rest of them contain length

zalando

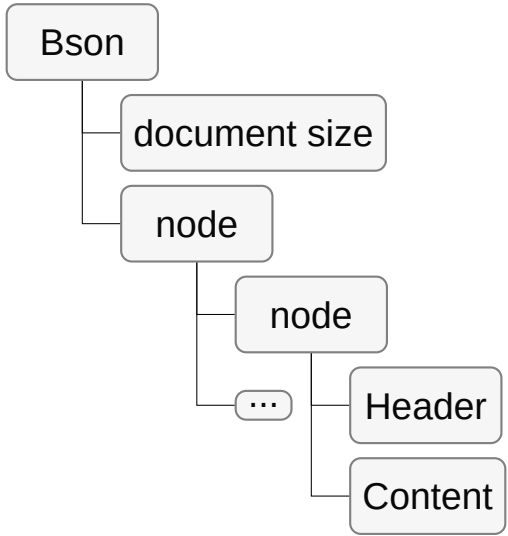

zalando

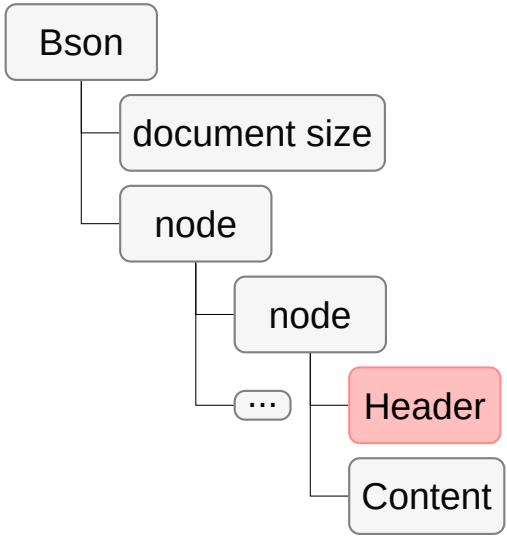

zalando

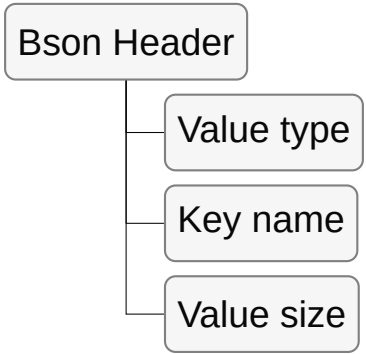

zalando

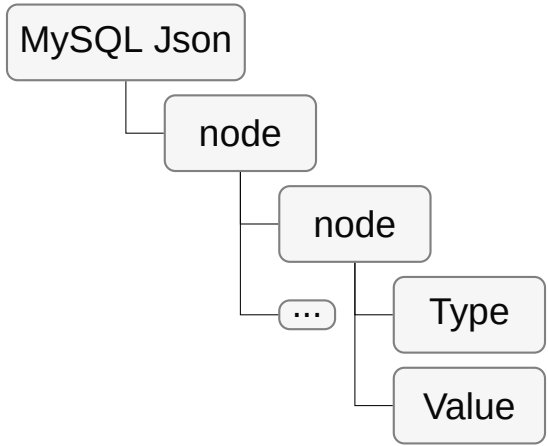

zalando

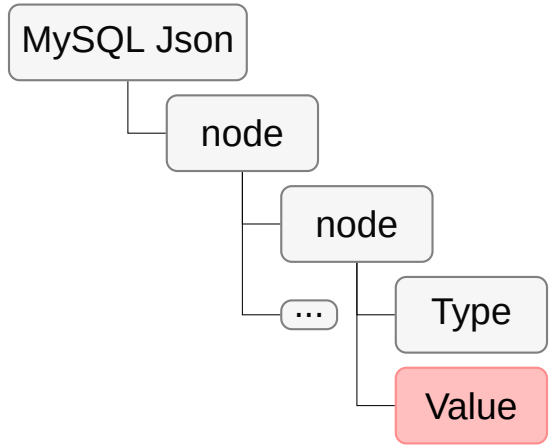

zalando

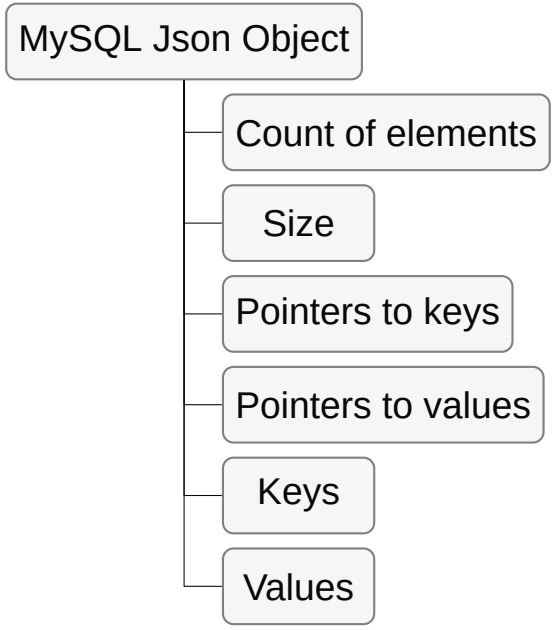
MySQL Json Object
Count of elements
Size
Pointers to keys
Pointers to values
Keys
Values

zalando

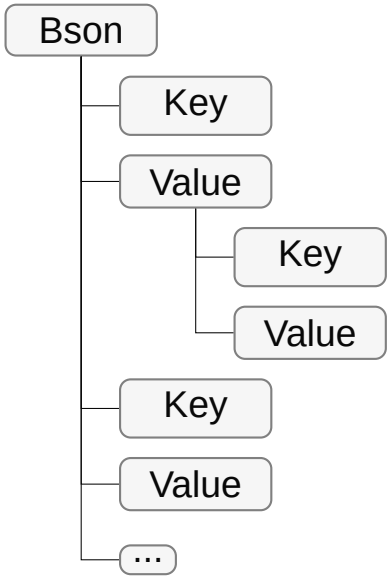
Bson
Key
Value
Key
Value
Key
Value
…

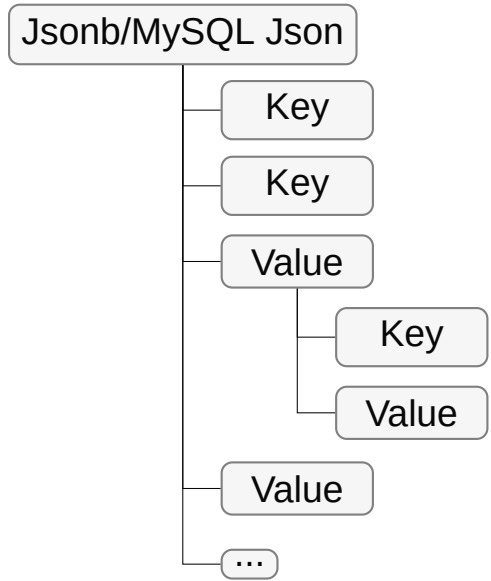
Jsonb/MySQL Json
Key
Key
Value
Key
Value
Value
…

zalando

```
{"a": 3, "b": "xyz"}
```


zalando

```
select pg_relation_filepath(oid),
relpages from pg_class
where relname = 'table_name';

 pg_relation_filepath | relpages
----------------------+----------
 base/40960/325477    |        0
(1 row)
```

```
\x10\x03\x00\x00\x00ab\x00\x00\x00\x00\x00\x00\x80\x03\x00xyz\x00\x00\x00\x00
```


zalando

```
bson.dumps({"a": 3, "b": u"xyz"})
```

```
\x17\x00\x00\x00\x10a\x00\x03\x00\x00\x00\x02b\x00\x04\x00\x00\x00xyz\x00\x00
```


zalando

$ hexdump -C database/table.ibd

```
\x00\x02\x00\x18\x00\x12\x00\x01\x00\x13\x00\x01\x00\x05\x03\x00\x0c\x14\x00ab\x03xyz\x00
```


zalando

TOAST

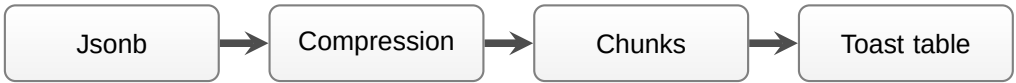

- ➔ TOAST_TUPLE_THRESHOLD bytes (normally 2 kB)
- ➔ PostgreSQL and MySQL use LZ variation
- ➔ MongoDB uses snappy block compression

Alignment

Variable-length portion is aligned to a 4-byte

```
insert into test
values('{"a": "aa", "b": 1}');
```

```
abaa\x20\x00\x00\x00\x00\x80\x01\x00
```

```
insert into test
values('{"a": 1, "b": "aa"}');
```

```
\x00\x00ab\x00\x00\x20\x00\x00\x00\x00\x80\x01\x00aa
```

In-memory representation

- ➔ Tree-like representation (JsonbValue, Document, Json_dom)
- ➔ Little bit more expensive but more convenient to work with
- ➔ Mostly in use to modify data (except MySQL)
- ➔ Most of the read operations use on-disk representation

Indexing support

- ➔ Postgresql – single path, multiple paths, entire document
- ➔ MongoDB – single path, multiple paths
- ➔ MySQL – virtual columns, single path, multiple paths

PG indexing details

- ➔ jsonb_path
- ➔ jsonb_path_ops

zalando

Queries

Pitfalls

- ➔ No Json path out of the box (jsquery, SQL/JSON)
- ➔ Queries with an array somewhere in the middle
- ➔ Iterating through document
- ➔ Update inside document

```
[{
    "items": [
        {"id": 1, "value": "aaa"},
        {"id": 2, "value": "bbb"}
    ]
}, {
    "items": [
        {"id": 3, "value": "aaa"},
        {"id": 4, "value": "bbb"}
    ]
}]
```


zalando

```
WITH items AS (
    SELECT jsonb_array_elements(data→'items')
    AS item FROM test
)
SELECT * FROM items
WHERE item-»'value' = 'aaa';
```

```
item
————————
 {"id": 1, "value": "aaa"}
 {"id": 3, "value": "aaa"}
(2 rows)
```


zalando

```
{
    "items": {
        "item1": {"status": true},
        "item2": {"status": true},
        "item3": {"status": false}
    }
}
```


zalando

```
WITH items AS (
    SELECT jsonb_each(data→'items')
    AS item FROM test
)
SELECT (item).key FROM items
WHERE (item).value-»'status' = 'true';
```

```
key
——-
item1
item2
(2 rows)
```


zalando

Benchmarks

GREAT PERFORMANCE

AWS EC2

m4.large instance

separate instance (database and generator)

16GB memory, 4 core 2.3GHz

Ubuntu 16.04

Same VPC and placement group

AMI that supports HVM virtualization type

at least 4 rounds of benchmark

PostgreSQL 9.6.3/10

MySQL 5.7.9/8.0

MongoDB 3.4.4

YCSB 0.13

10^6 rows and operations

AWS EC2

Configuration

shared_buffers
effective_cache_size
max_wal_size
innodb_buffer_pool_size
innodb_log_file_size
write concern level (journaled or transaction_sync)
checkpoint
eviction

zalando

Document types

“simple” document
10 key/value pairs (100 characters)

“large” document
100 key/value pairs (200 characters)

“complex” document
100 keys, 3 nesting levels (100 characters)

zalando

Select, GIN

”simple” document

jsonb_path_ops

zalando

Throughput (ops/sec)

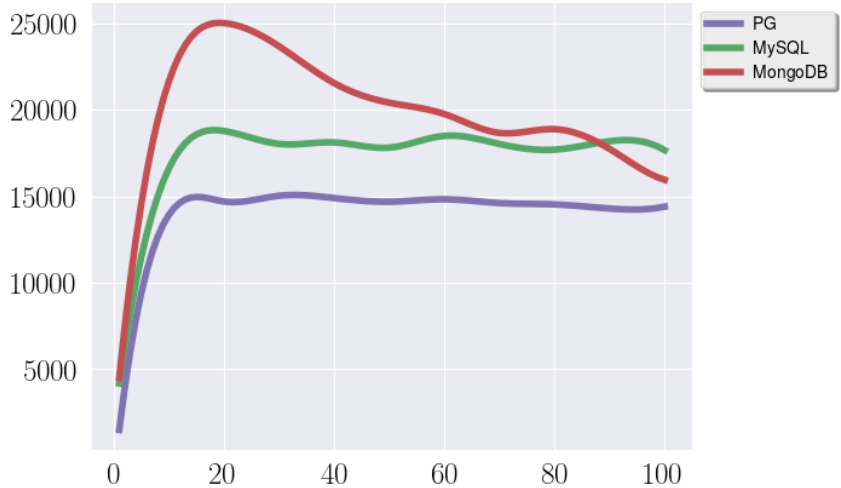

zalando

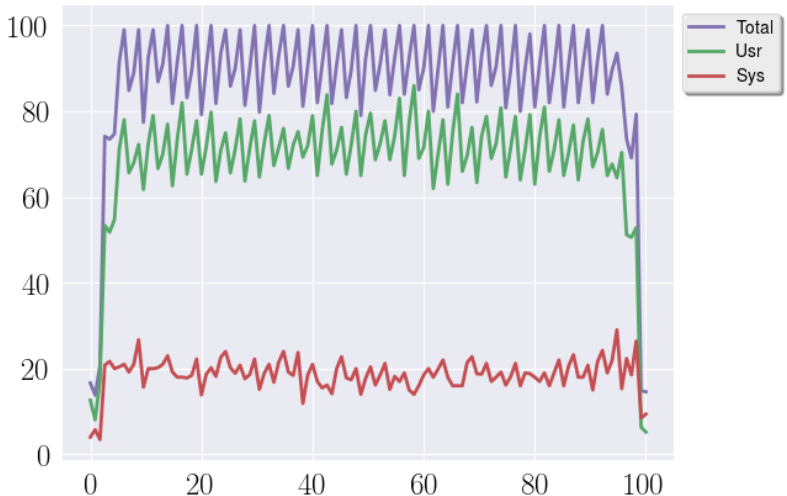

zalando

CPU migrations (MongoDB)

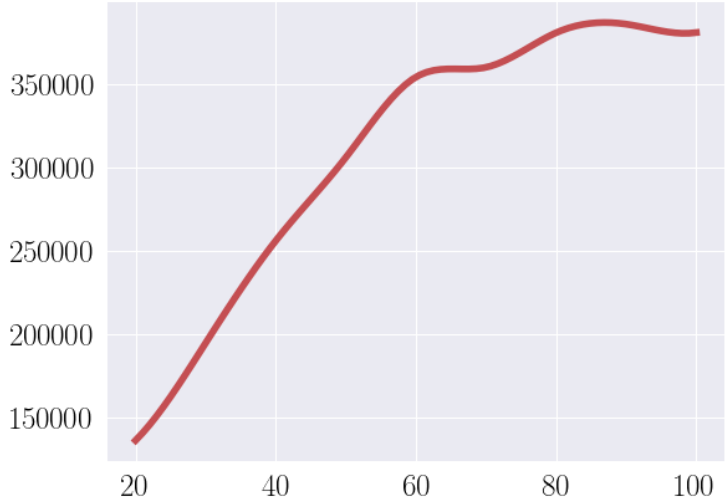

zalando

Latency 99% (μs)

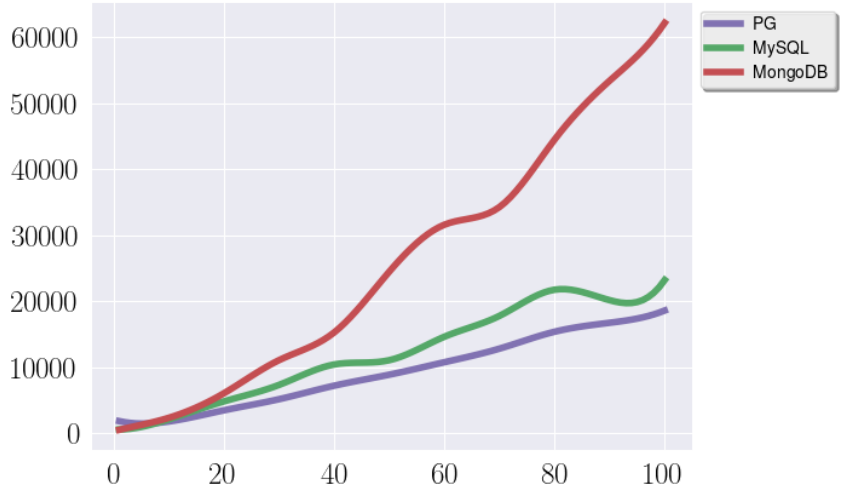

zalando

Select, BTree

”simple” document

btree

zalando

Throughput (ops/sec)

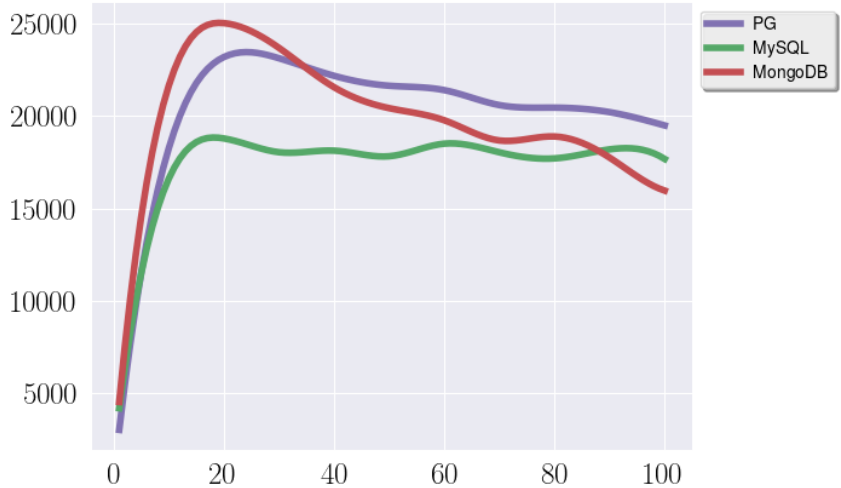

zalando

Select, BTree

”complex” document

btree

zalando

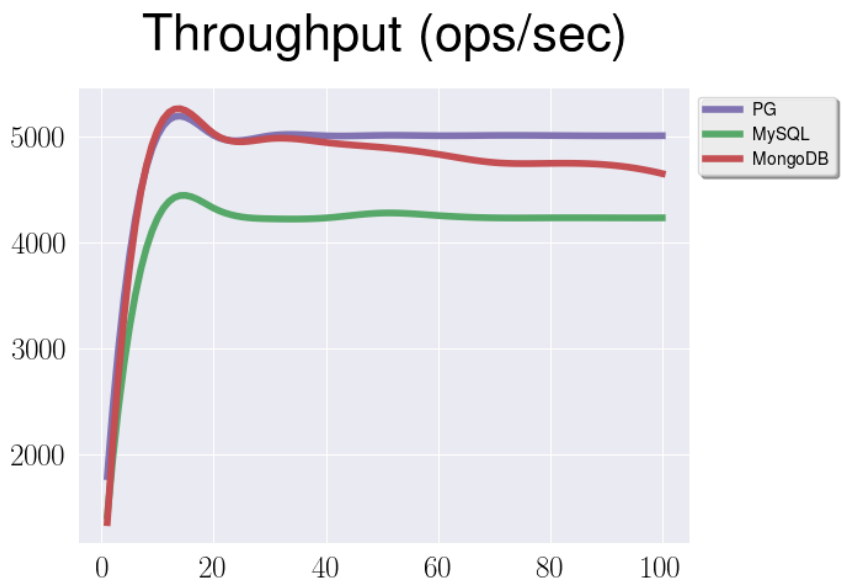
Throughput (ops/sec)
5000
4000
3000
2000
0
20
40
60
80
100
PG
MySQL
MongoDB

zalando

Insert

”simple” document

journaled

zalando

Throughput (ops/sec)

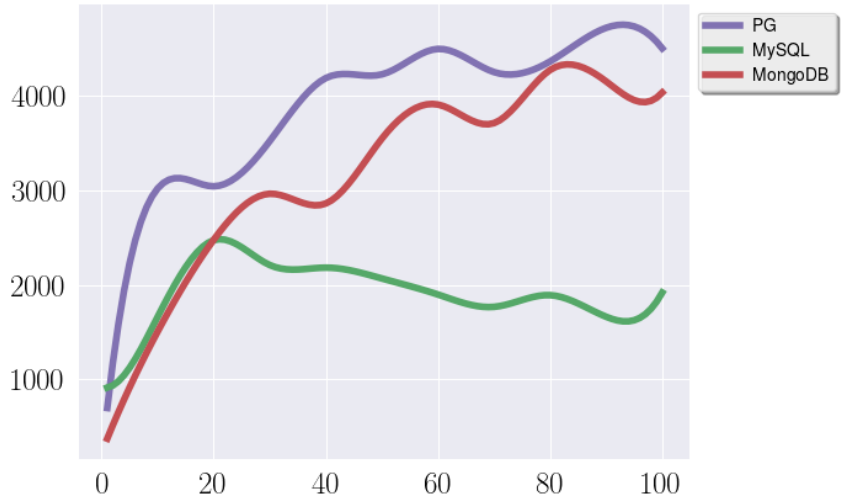

zalando

CPU%

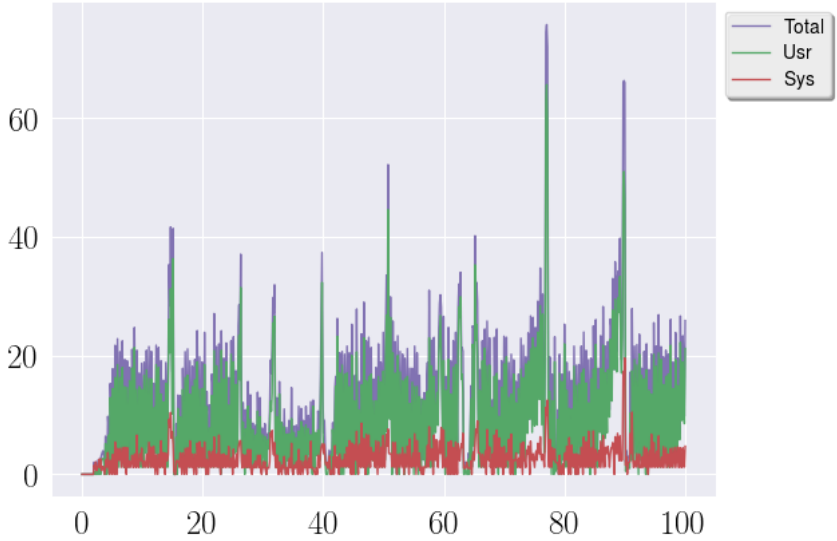

zalando

IO queue size

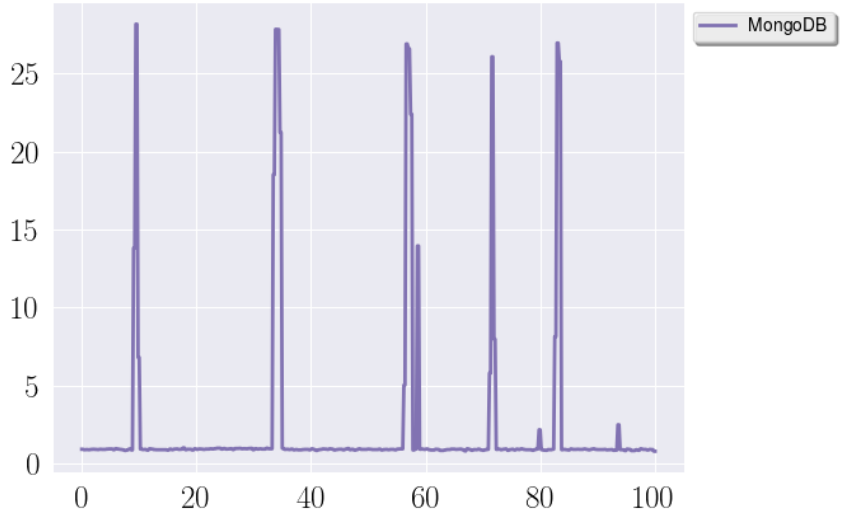

zalando

Throughput (ops/sec)

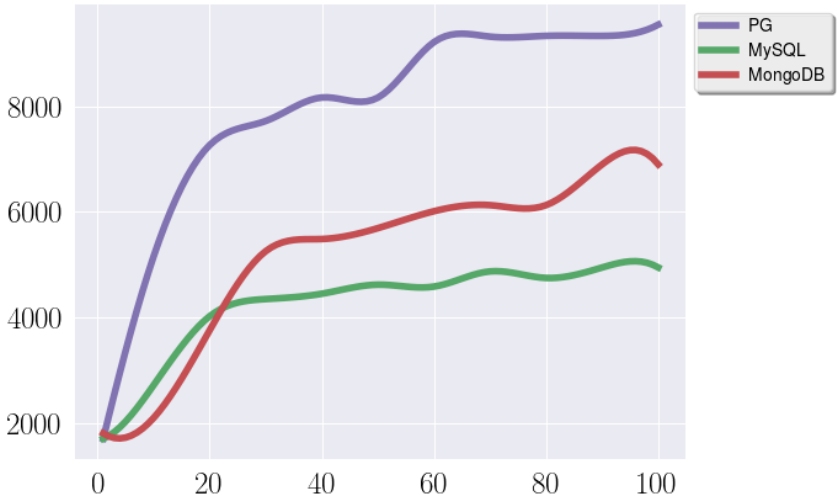

zalando

Update 50%, Select 50%

”simple” document

Update one field

journaled

max wal size 4GB

zalando

Throughput (ops/sec)

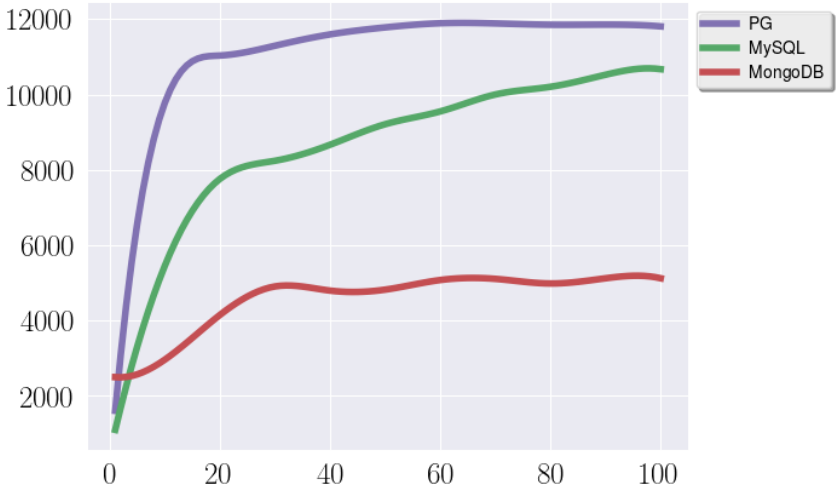

zalando

Update 50%, Select 50%

”large” document

Update one field

zalando

Throughput (ops/sec)

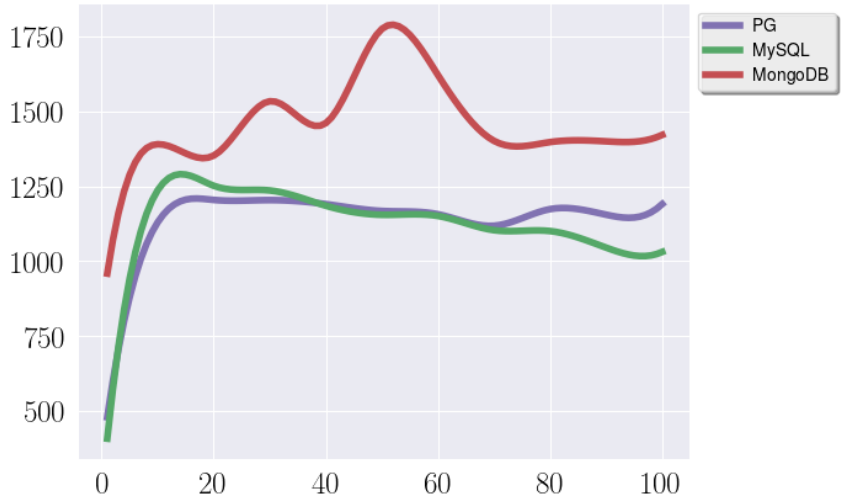

zalando

JSON vs JSONB

”simple” document

btree

insert

zalando

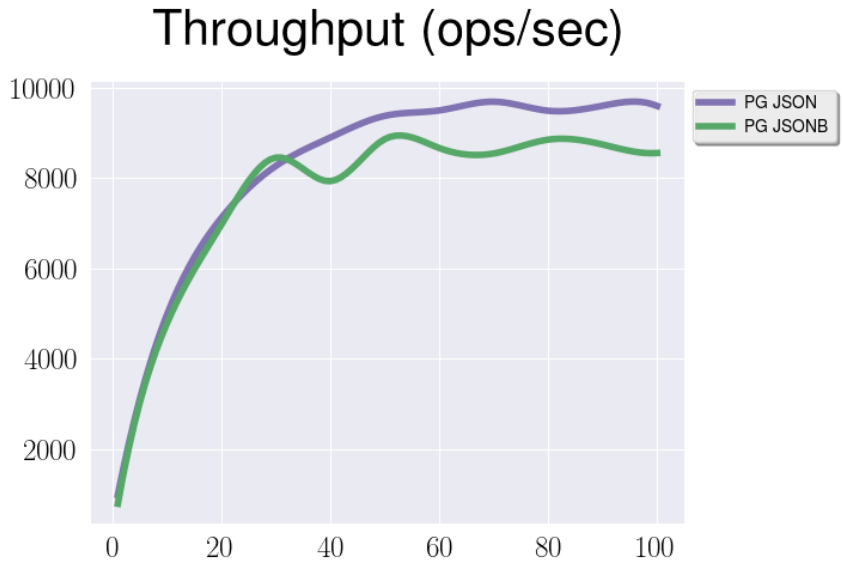
Throughput (ops/sec)
10000
8000
6000
4000
2000
0
20
40
60
80
100
PG JSON
PG JSONB

zalando

JSON vs JSONB

”simple” document

btree

select

zalando

Throughput (ops/sec)

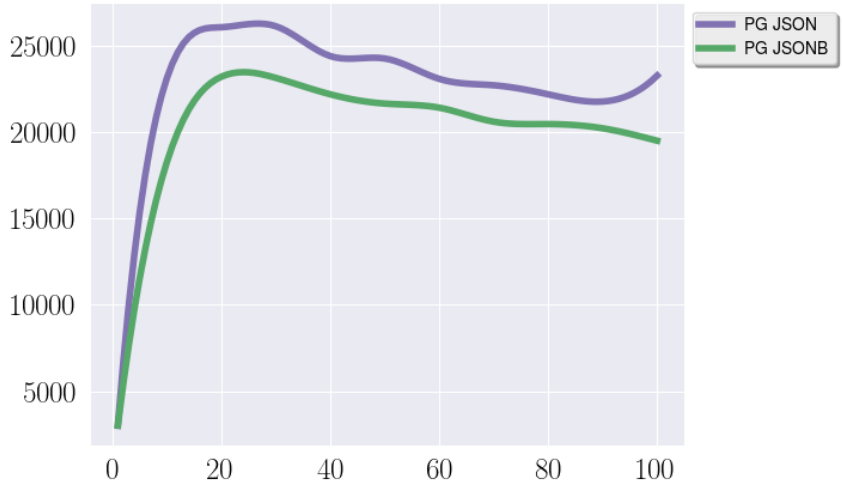

zalando

SQL vs JSONB

”simple” document

btree

insert

zalando

Throughput (ops/sec)

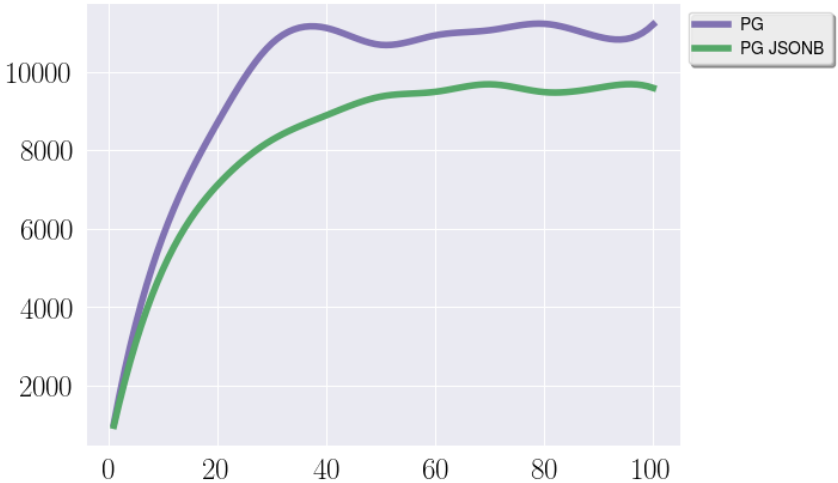

zalando

SQL vs JSONB

”simple” document

btree

select

zalando

Throughput (ops/sec)
25000
20000
15000
10000
5000
0
20
40
60
80
100
PG
PG JSONB

zalando

How to bring it down acci-
dentally?

zalando

- ➔ Update one field of a document
- ➔ DETOAST of a document
  (select, constraints, procedures etc.)
- ➔ Reindex of an entire document

zalando

```
select data→'key1'→'key2' from table;

select data→'key1', data→'key2' from table;
```


zalando

Document slice

”large” document

One field from a document

zalando

Throughput (ops/sec)

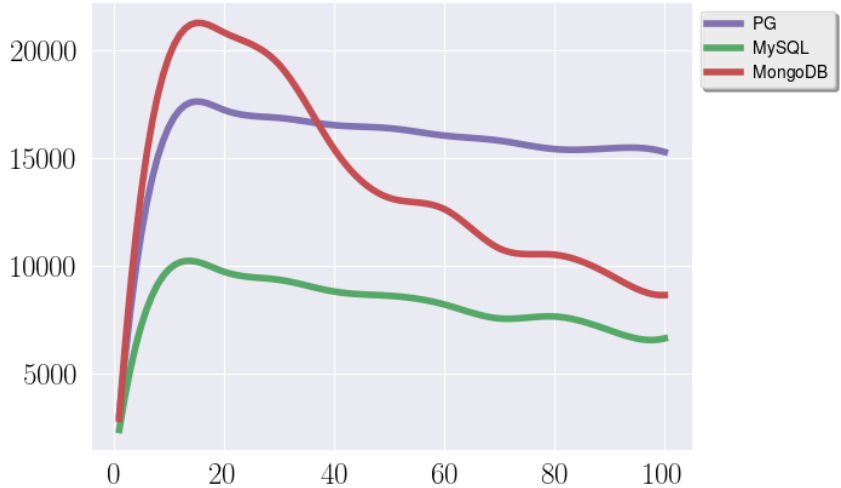

zalando

Document slice

”large” document

10 fields from a document

zalando

Throughput (ops/sec)

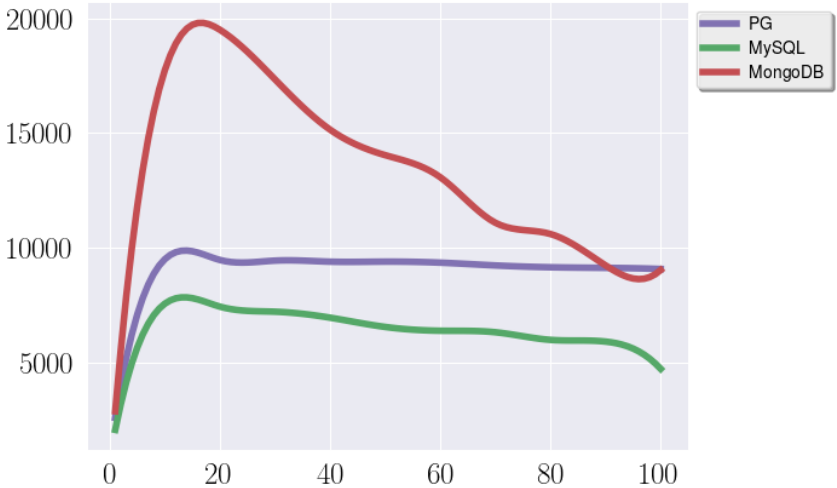

zalando

Document slice

```
create type test as ("a" text, "b" text);
insert into test_jsonb
values('{"a": 1, "b": 2, "c": 3}');
select q.* from test_jsonb,
jsonb_populate_record(NULL::test, data) as q;
 a | b
---+---
 1 | 2
(1 row)
```


zalando

SET STORAGE EXTERNAL

TOAST_TUPLE_THRESHOLD

”simple” document

40 threads

different document size

select

zalando

Throughput, 40 clients

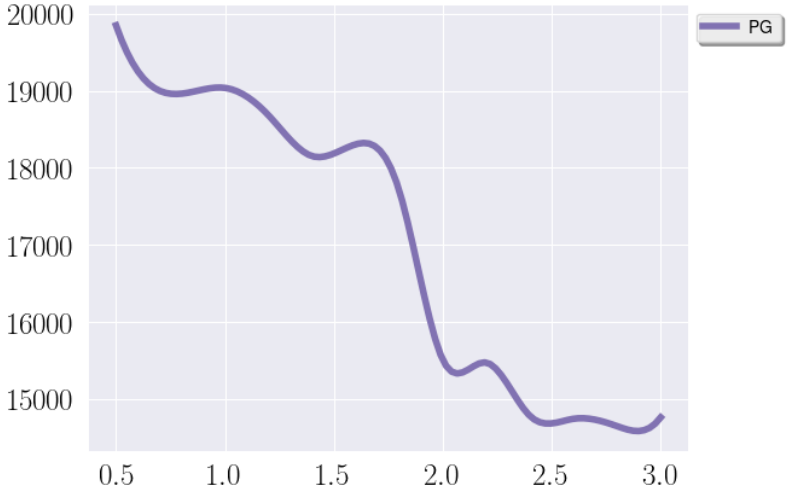

zalando

Index update

”simple” document

Update one field

zalando

PG, Throughput (ops/sec)

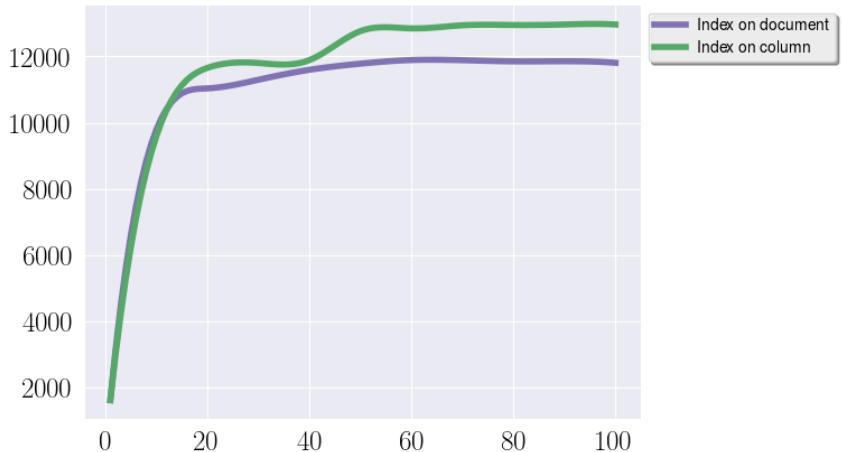

zalando

➔ Jsonb is more that good for many use cases

zalando

- ➔ Jsonb is more that good for many use cases
- ➔ Benchmarks above are only ”hints”

zalando

- ➔ Jsonb is more that good for many use cases
- ➔ Benchmarks above are only ”hints”
- ➔ You need your own tests

zalando

Questions?

github.com/erthalion

@erthalion

9erthalion6 at gmail dot com

zalando