



NOSQL BEST PRACTICES

FOR POSTGRESQL



DMITRY DOLGOV

29-03-2018



Introduction

Less benchmarks

More opinionated best practices

Introduction

Application developers

DBAs

Extension developers

Application developers

When jsonb?

When jsonb?

→ You have a distinct flexible model

When jsonb?

- You have a distinct flexible model
- You need to work with data provided in document oriented format

When jsonb?

- You have a distinct flexible model
- You need to work with data provided in document oriented format
- Workaround for technical issues (large number of tables or expensive alignment)

When not jsonb?

When not jsonb?

→ Flexibility "just in case"


When not jsonb?

- Flexibility "just in case"
- Reluctance to create a migration

When not jsonb?

- Flexibility "just in case"
- Reluctance to create a migration
- Use jsonb column as a "garbage can"

Jsonb -> Relation



When to move from jsonb to relation?

When to move from jsonb to relation?

- Queries rely significantly in information about internal structure of documents

When to move from jsonb to relation?

- Queries rely significantly in information about internal structure of documents
- There are too many constraints for documents

When to move from jsonb to relation?

- Queries rely significantly in information about internal structure of documents
- There are too many constraints for documents
- Some parts of document are used much more frequently than other

```
SELECT id, created FROM some_table
WHERE
```

```
(data->>'name' = :a
AND (data @> ('{"items":[{"id":"' || :b || '"}]}'))
AND (data @> ('{"items":[{"elems":[{"name":"' || :c || '"}]}]}'))
AND (data @> ('{"items":[{"elems":[{"id":"' || :d || '"}]}]}'))
AND (data @> ('{"items":[{"name":"' || :e || '"}]}'))

ORDER BY created ASC, id ASC;
```

Complicated conditions

- jquery
- SQL/JSON

Complicated conditions

```
SELECT id, created FROM some_table  
WHERE
```

```
data @ 'items.#{id = '||:a||'}'  
AND data @ 'items.#{elems.#{name = '||:b||'}'  
AND data @ 'items.#{elems.#{id = '||:c||'}'  
AND data @ 'items.#{name = '||:d||'}'
```

```
ORDER BY created ASC, id ASC;
```

Complicated conditions

```
SELECT id, created FROM some_table  
WHERE
```

```
data @~ '$.items[*] ? (@id = '||:a||')'  
AND data @~ '$.items[*].elems[*] ? (@name = '||:b||')'  
AND data @~ '$.items[*].elems[*] ? (@id = '||:c||')'  
AND data @~ '$.items[*](@name = '||:d||')
```

```
ORDER BY created ASC, id ASC;
```


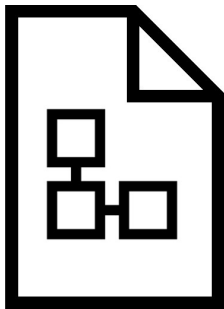
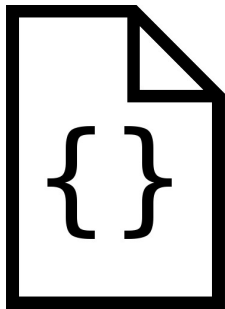
Complicated select

SELECT

```
st.data #>> '{item_a, another_item}' AS item_a,  
st.data #>> '{item_c}' AS item_c,  
jsonb_array_elements(  
    data #> '{item_b, subitem_a, subitem_b}'  
) ->> 'some_key' AS item_e
```

```
FROM some_table st LEFT JOIN another_table at  
ON (st.data #> '{item_b, key_a, key_b}') @>  
    jsonb_build_array(jsonb_build_object(  
        'key', 'some_key_name',  
        'value', at.data #>> '{item_b, another_item}'  
    ));
```

Jsonb -> Relation



Constraints

- Simple checks for value, type or size
- More convenient checks with jquery
- Json schema

Constraints

- Simple checks for value, type or size
- More convenient checks with jquery
- Json schema

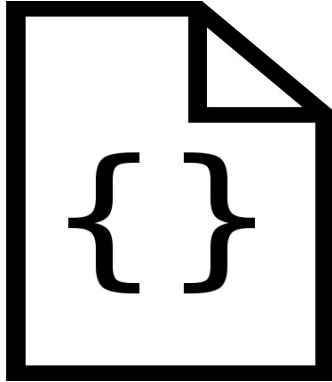
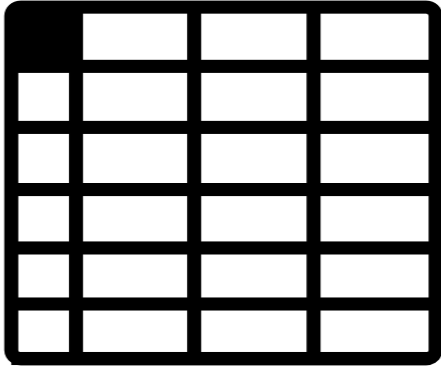
Constraints

```
CREATE TABLE test (  
  data jsonb,  
  CHECK (jsonb_typeof(data->'key') = 'array')  
);
```

```
CREATE TABLE test (  
  data jsonb,  
  CHECK (data @> 'key IS ARRAY OR key IS OBJECT')  
);
```

```
CREATE TABLE test (  
  data jsonb,  
  CHECK (validate_json_schema('{"key": "array"}', data))  
);
```

Relation -> Jsonb



```
SELECT jsonb_agg(query) FROM (  
    SELECT id, data  
    FROM jsonb_table  
) query;
```

Seamless interaction between json and rela- tion

```
[{
  "items": [
    {"id": 1, "value": "aaa"},
    {"id": 2, "value": "bbb"}
  ]
}, {
  "items": [
    {"id": 3, "value": "aaa"},
    {"id": 4, "value": "bbb"}
  ]
}]
```

```
WITH items AS (  
    SELECT jsonb_array_elements(data->'items')  
    AS item FROM test  
)  
SELECT * FROM items  
WHERE item->>'value' = 'aaa';  
  
item  
-----  
{ "id": 1, "value": "aaa" }  
{ "id": 3, "value": "aaa" }  
(2 rows)
```

```
WITH items AS (  
    SELECT jsonb_array_elements(data->'items')  
    AS item FROM test  
)  
SELECT * FROM items  
WHERE item->>'value' = 'aaa';  
  
item  
-----  
{ "id": 1, "value": "aaa" }  
{ "id": 3, "value": "aaa" }  
(2 rows)
```



```
{  
  "items": {  
    "item1": {"status": true},  
    "item2": {"status": true},  
    "item3": {"status": false}  
  }  
}
```

```
WITH items AS (  
    SELECT jsonb_each(data->'items')  
    AS item FROM test  
)  
SELECT (item).key FROM items  
WHERE (item).value->>'status' = 'true';
```

key

item1

item2

(2 rows)

```
WITH items AS (  
    SELECT jsonb_each(data->'items')  
    AS item FROM test  
)  
SELECT (item).key FROM items  
WHERE (item).value->>'status' = 'true';
```

key

item1

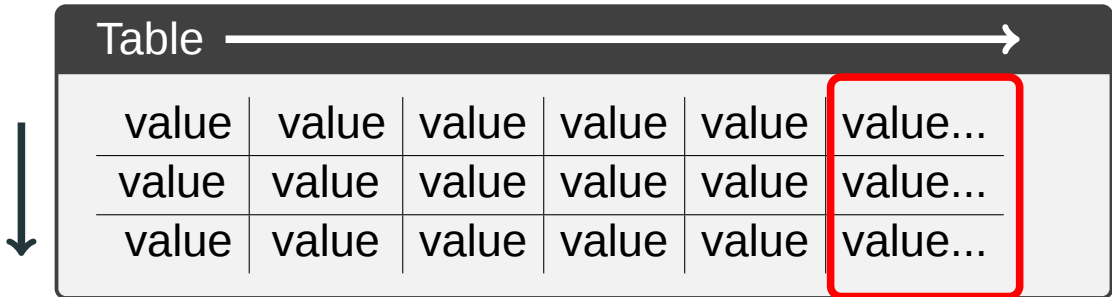
item2

(2 rows)

Multiple jsonb columns

- Keep at the end for readability
- tuple_deform

Multiple jsonb columns



The diagram illustrates a table structure with multiple jsonb columns. A dark grey header bar at the top is labeled "Table" on the left and features a white arrow pointing to the right. Below the header, a table with a light grey background is shown. It consists of 6 columns and 3 rows. The first five columns each contain the word "value". The sixth column contains "value...", "value...", and "value..." in its three rows respectively. This sixth column is highlighted with a red rectangular border. To the left of the table, a dark grey arrow points downwards, indicating the rows. The word "Table" is positioned to the left of the header bar.

Table					
value	value	value	value	value	value...
value	value	value	value	value	value...
value	value	value	value	value	value...

Multiple jsonb columns

Table →

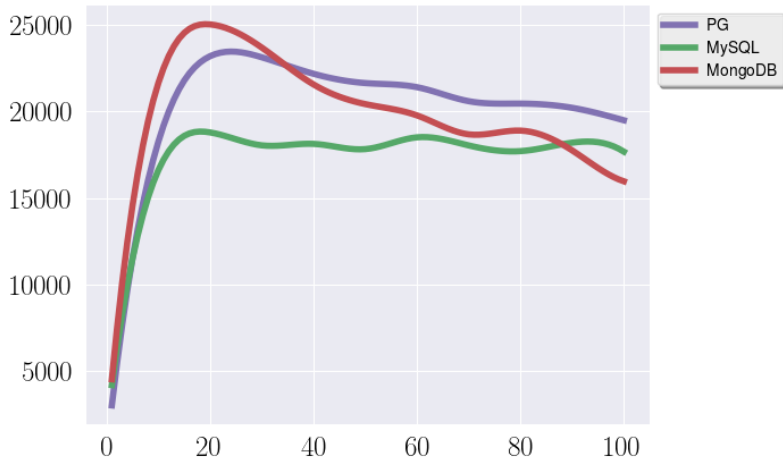
↓

value...	value	value	value	value	value
value...	value	value	value	value	value
value...	value	value	value	value	value


Document slice: in the DB or on the app?

- Amount of data passed from DB to application
- Performance hit in some cases
(multiple detoasting)

Throughput (ops/sec)




Throughput (ops/sec)




- Plain Json
- Binary Jsonb
- Relation

Throughput (ops/sec)



Throughput (ops/sec)



Jsonb array vs regular array

- Store elements of different type?
Not really a "single model" idea.
- Syntax is less natural (this may change)
- Updates are slower
- Arrays are 1-based, Jsonb 0-based

```
SELECT array[0] FROM some_table;
SELECT jsonb->0 FROM some_table;
-- WIP
SELECT jsonb[0] FROM some_table;

UPDATE some_table SET array[0] = 'new_value';
UPDATE some_table
SET jsonb = jsonb_set(jsonb, '{0}', 'new_value');
-- WIP
UPDATE some_table SET jsonb[0] 'new_value';
```

Some useful extensions

- jsquery
- postgres-json-schema
- is_jsonb_valid
- zson (jsonbc, custom compression methods WIP)
- jsonb_explorer

Types, please



DBAs

Limitations

Size 256 MB

Depth - max_stack_depth

Stack depth is different for create & update

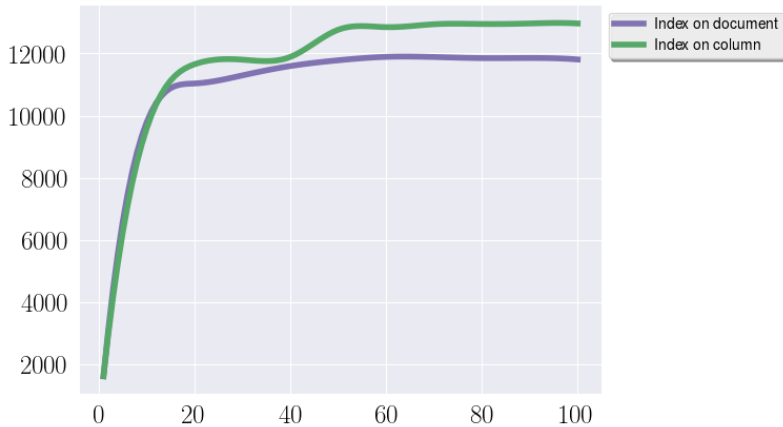
Indexing support

- GIN index (jsonb_ops, jsonb_path_ops)
- Functional BTree index
- jsquery strategies for GIN
- Partial indexes WIP

Place for ID

- Inside a document
- As a separate column

PG, Throughput (ops/sec)



Place for ID

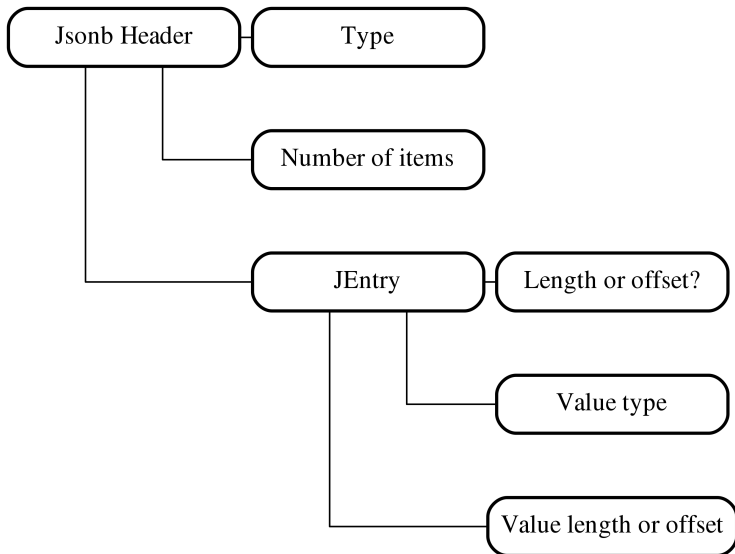
PostgreSQL 11 have HOT updates for some expression indexes, which will eliminate this problem.


Statistics

- There is no proper selectivity estimation for jsonb
- Optimizer can give wrong estimations for GIN and complex queries
- Functional indexes

How much to write?

Jsonb vs Json






How much to write?

- Every update leads to update of an entire document (but it's ok)
- WAL can have a full document or just a diff
- Old and new tuples fit into the same page - diff
- Old and new tuples fit into the same page - full
- If logical decoding is enabled - full


How much to write?



Huge documents

- TOAST has significant overhead (assemble, locks)
- Other than that linear degradation

Throughput, 40 clients



Alignment

Variable-length portion is aligned to a 4-byte

```
insert into test  
values('{"a": "aa", "b": 1}');
```

```
abaa\x20\x00\x00\x00\x00\x80\x01\x00
```

```
insert into test  
values('{"a": 1, "b": "aa"}');
```

```
\x00\x00ab\x00\x00\x20\x00\x00\x00\x00\x80\x01\x00aa
```


Extensions

Why to write an extension?

- Implement some convenient functionality (e.g. jsonb intersection)
- Create function optimized for your domain model

Why to write an extension?

findJsonValueFromContainer -> binary search
_id in a fixed position?

- Raw Jsonb container when search for an element
- Iterate through JsonbValue when update

Reuse infrastructure

- findJsonValueFromContainer
- JsonbIterator
- addToParseState
- worker functions


Random tips

- Clone iterator
- String are not null-terminated

Questions?

 github.com/erthalion

 [@erthalion](https://twitter.com/erthalion)

 9erthalion6 at gmail dot com