



# NOSQL BEST PRACTICES

---

FOR POSTGRESQL



DMITRY DOLGOV

20-04-2018



# Introduction

Less benchmarks


More opinionated best practices

# Introduction

Application developers

DBAs

Extension developers



# Application developers

# When jsonb?

## When jsonb?

→ You have a distinct flexible model

## When jsonb?

- You have a distinct flexible model
- You need to work with data provided in document oriented format

## When jsonb?

- You have a distinct flexible model
- You need to work with data provided in document oriented format
- Workaround for technical issues (large number of tables or expensive alignment)



## When not jsonb?

## When not jsonb?

→ Flexibility "just in case"

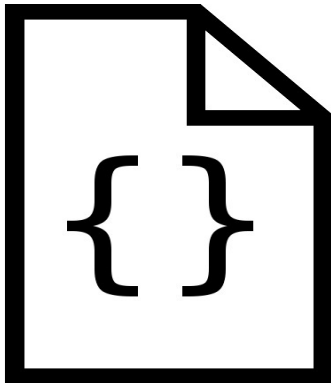
## When not jsonb?

- Flexibility "just in case"
- Reluctance to create a migration

## When not jsonb?

- Flexibility "just in case"
- Reluctance to create a migration
- Use jsonb column as a "garbage can"

## Jsonb -> Relation



# When to move from jsonb to relation?

## When to move from jsonb to relation?

- Queries rely significantly in information about internal structure of documents

## When to move from jsonb to relation?

- Queries rely significantly in information about internal structure of documents
- There are too many constraints for documents



## When to move from jsonb to relation?

- Queries rely significantly in information about internal structure of documents
- There are too many constraints for documents
- Some parts of document are used much more frequently than other

```
SELECT id, created FROM some_table  
WHERE
```

```
(data->'name' = :a  
AND (data @> ('{"items":[{"id":"" || :b || ""}]}'))  
AND (data @> ('{"items":[{"elems":[{"name":"" || :c || ""}]}]}'))  
AND (data @> ('{"items":[{"elems":[{"id":"" || :d || ""}]}]}'))  
AND (data @> ('{"items":[{"name":"" || :e || ""}]}'))  
  
ORDER BY created ASC, id ASC;
```

## Complicated conditions

- jquery
- SQL/JSON

## Complicated conditions

```
SELECT id, created FROM some_table  
WHERE
```

```
  data @@ 'items.#{id = ' || :a || '}'  
AND data @@ 'items.#{elems.#{name = ' || :b || '}'  
AND data @@ 'items.#{elems.#{id = ' || :c || '}'  
AND data @@ 'items.#{name = ' || :d || '}'
```

```
ORDER BY created ASC, id ASC;
```

## Complicated conditions

```
SELECT id, created FROM some_table  
WHERE
```

```
data @~ '$.items[*] ? (@id = ' || :a || ' )'  
AND data @~ '$.items[*].elems[*] ? (@name = ' || :b || ' )'  
AND data @~ '$.items[*].elems[*] ? (@id = ' || :c || ' )'  
AND data @~ '$.items[*](@name = ' || :d || ' )'
```

```
ORDER BY created ASC, id ASC;
```

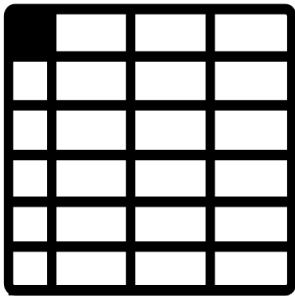
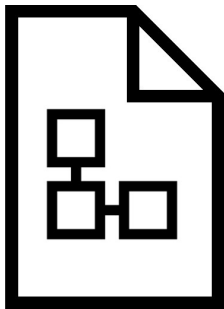
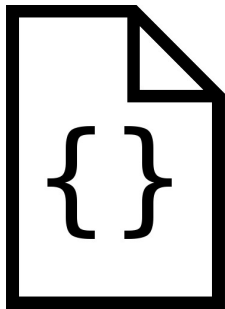
# Complicated select

SELECT

```
st.data #» '{item_a, another_item}' AS item_a,  
st.data #» '{item_c}' AS item_c,  
jsonb_array_elements(  
    data #> '{item_b, subitem_a, subitem_b}'  
) -» 'some_key' AS item_e
```

```
FROM some_table st LEFT JOIN another_table at  
ON (st.data #> '{item_b, key_a, key_b}') @>  
    jsonb_build_array(jsonb_build_object(  
        'key', 'some_key_name',  
        'value', at.data #» '{item_b, another_item}'  
    ));
```

## Jsonb -> Relation



## Constraints

- Simple checks for value, type or size
- More convenient checks with jquery
- Json schema



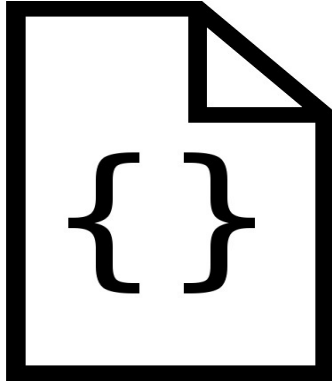
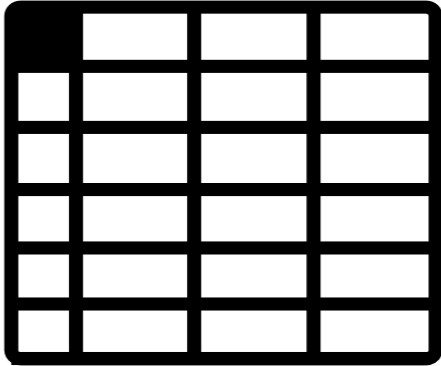
# Constraints

```
CREATE TABLE test (  
  data jsonb,  
  CHECK (jsonb_typeof(data→'key') = 'array')  
);
```

```
CREATE TABLE test (  
  data jsonb,  
  CHECK (data @@ 'key IS ARRAY OR key IS OBJECT')  
);
```

```
CREATE TABLE test (  
  data jsonb,  
  CHECK (validate_json_schema('{ "key": "array" }', data))  
);
```

## Relation -> Jsonb



```
SELECT jsonb_agg(query) FROM (  
    SELECT id, data  
    FROM jsonb_table  
) query;
```

# Seamless interaction between json and relation

```
[{
  "items": [
    {"id": 1, "value": "aaa"},
    {"id": 2, "value": "bbb"}
  ]
}, {
  "items": [
    {"id": 3, "value": "aaa"},
    {"id": 4, "value": "bbb"}
  ]
}]
```

```
WITH items AS (  
    SELECT jsonb_array_elements(data->'items')  
    AS item FROM test  
)  
SELECT * FROM items  
WHERE item->>'value' = 'aaa';  
  
item  
-----  
{"id": 1, "value": "aaa"}  
{"id": 3, "value": "aaa"}  
(2 rows)
```

```
WITH items AS (  
    SELECT jsonb_array_elements(data->'items')  
    AS item FROM test  
)  
SELECT * FROM items  
WHERE item->>'value' = 'aaa';  
  
item  
-----  
{"id": 1, "value": "aaa"}  
{"id": 3, "value": "aaa"}  
(2 rows)
```

```
{  
  "items": {  
    "item1": {"status": true},  
    "item2": {"status": true},  
    "item3": {"status": false}  
  }  
}
```



```
WITH items AS (  
    SELECT jsonb_each(data->'items')  
    AS item FROM test  
)  
SELECT (item).key FROM items  
WHERE (item).value->>'status' = 'true';  
  
key  
--  
item1  
item2  
(2 rows)
```

```
WITH items AS (  
    SELECT jsonb_each(data->'items')  
    AS item FROM test  
)  
SELECT (item).key FROM items  
WHERE (item).value->>'status' = 'true';  
  
key  
--  
item1  
item2  
(2 rows)
```

## Multiple jsonb columns

- Keep at the end for readability
- tuple\_deform (PG11, JIT compilation)

## Multiple jsonb columns

Table

value	value	value	value	value	value...
value	value	value	value	value	value...
value	value	value	value	value	value...

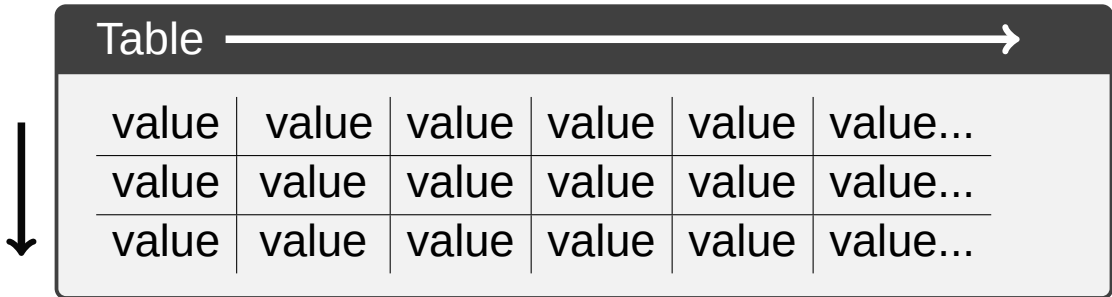
## Multiple jsonb columns

Table



value	value	value	value	value	value...
value	value	value	value	value	value...
value	value	value	value	value	value...

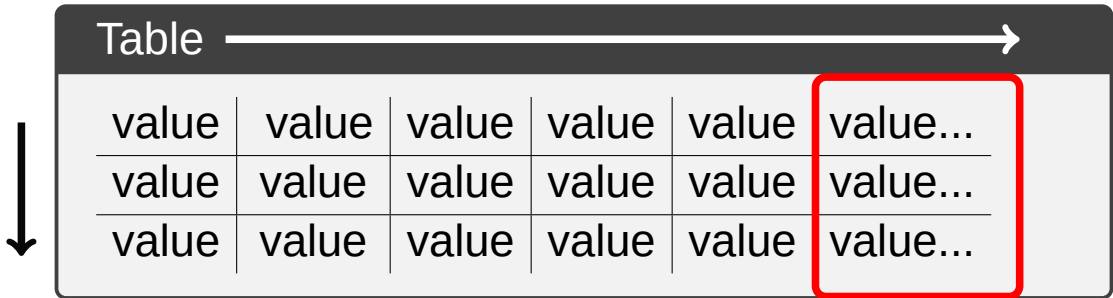
# Multiple jsonb columns



The diagram illustrates a table structure with multiple jsonb columns. A dark grey header bar at the top is labeled "Table" on the left and features a white arrow pointing to the right. Below this header is a light grey table area. To the left of the table, a black arrow points downwards, indicating the rows. The table itself consists of three rows and six columns. Each cell in the table contains the text "value" followed by an ellipsis ("value...").

value...	value...	value...	value...	value...	value...
value...	value...	value...	value...	value...	value...
value...	value...	value...	value...	value...	value...

## Multiple jsonb columns



The diagram illustrates a table structure with multiple jsonb columns. A dark grey header bar at the top is labeled "Table" on the left and features a white arrow pointing to the right. Below the header, a light grey table contains six columns. The first five columns each contain the word "value" in all three rows. The sixth column contains "value..." in all three rows and is highlighted with a red rectangular border. To the left of the table, a black arrow points downwards, indicating the rows. The table structure is as follows:

Table →					
value	value	value	value	value	value...
value	value	value	value	value	value...
value	value	value	value	value	value...

## Multiple jsonb columns

Table →

↓

value...	value	value	value	value	value
value...	value	value	value	value	value
value...	value	value	value	value	value



## **Document slice: in the DB or in the app?**

- Amount of data passed from DB to application
- Performance hit in some cases  
(multiple detoasting)

```
select data->'key1'->'key2' from table;
```

```
select data->'key1', data->'key2' from table;
```


```
select data->'key1' -> 'key2' from table;
```

```
select data->'key1', data->'key2' from table;
```

```
select data->'key1'->'key2' from table;
```

```
select data->'key1', data->'key2' from table;
```

# Read latency, ms



- Plain Json
- Binary Jsonb
- Relation


# Insert workload

# Throughput (ops/sec)






# Throughput (ops/sec)




# Read workload

# Throughput (ops/sec)



# Throughput (ops/sec)



## Jsonb array vs regular array

- Store elements of different type?  
Not really a "single model" idea.
- Syntax is less natural (this may change)
- Updates are slower
- Arrays are 1-based, Jsonb 0-based

```
SELECT array[0] FROM some_table;
```

```
SELECT jsonb→0 FROM some_table;
```

– WIP

```
SELECT jsonb[0] FROM some_table;
```

```
UPDATE some_table SET array[0] = 'new_value';
```

```
UPDATE some_table
```

```
SET jsonb = jsonb_set(jsonb, '{0}', 'new_value');
```

– WIP

```
UPDATE some_table SET jsonb[0] = 'new_value';
```

## Jsonb NULL != SQL NULL

```
SELECT jsonb_set(data, '{key}', NULL);
```

```
jsonb_set
```

```
-----  
NULL
```

```
(1 row)
```

```
SELECT jsonb_set(data, '{key}', 'null');
```

```
jsonb_set
```

```
-----  
{"key": null}
```

```
(1 row)
```

## Some useful extensions

- jsquery
- postgres-json-schema
- is\_jsonb\_valid
- zson (custom compression methods WIP)
- jsonb\_explorer



**Types, please**



# DBAs

## Limitations

Size 256 MB

Depth - max\_stack\_depth

Stack depth is different for create & update




## Indexing support

- GIN index (jsonb\_ops, jsonb\_path\_ops)
- Functional BTree index
- jsquery strategies for GIN
- Partial indexes WIP

## Place for ID

- Inside a document
- As a separate column

## PG, Throughput (ops/sec)





## Place for ID

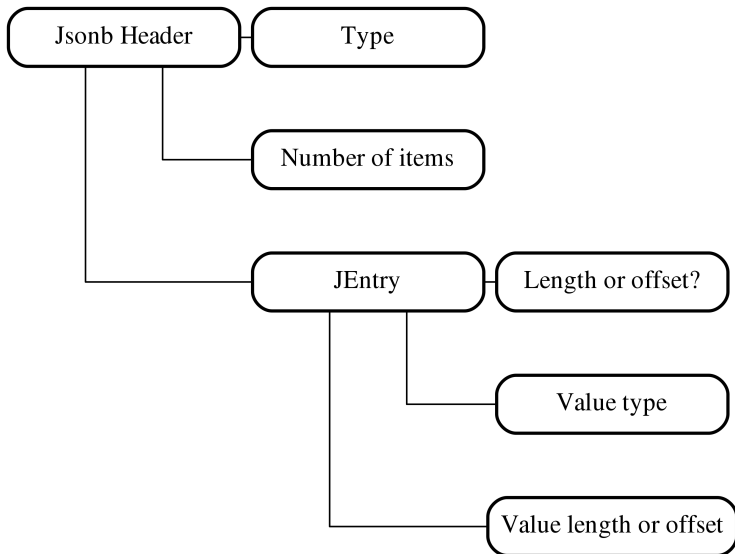
PostgreSQL 11 have HOT updates for some expression indexes, which will eliminate this problem.


## Statistics

- There is no proper selectivity estimation for jsonb
- Optimizer can give wrong estimations for GIN and complex queries
- Functional indexes

How much to write?

# Jsonb vs Json







## How much to write?

- Every update leads to update of an entire document (but it's ok)
- WAL can have a full document or just a diff
- Old and new tuples fit into the same page - diff
- If logical decoding is enabled - full

## How much to write?




## How much to write?







## How much to write?



## How much to write?




## How much to write?



## Huge documents

- TOAST has significant overhead (assemble, locks)
- Other than that linear degradation

## Throughput, 40 clients



# Alignment

Variable-length portion is aligned to a 4-byte

```
insert into test  
values( '{"a": "aa", "b": 1}' );
```

```
abaa\x20\x00\x00\x00\x00\x80\x01\x00
```

```
insert into test  
values( '{"a": 1, "b": "aa"}' );
```

```
\x00\x00ab\x00\x00\x20\x00\x00\x00\x00\x80\x01\x00aa
```

# Extensions

## Why write an extension?

- Implement some convenient functionality (e.g. jsonb intersection)
- Create function optimized for your domain model




## Why write an extension?

{	<b>_id</b>				}
---	------------	--	--	--	---


## Why write an extension?



## Why write an extension?



## Why write an extension?



- Raw Jsonb container when search for an element
- Iterate through JsonbValue when update

## Reuse infrastructure

- findJsonValueFromContainer
- JsonbIterator
- addToParseState
- worker functions


## Random tips


- Clone iterator
- String are not null-terminated

## Questions?

 [github.com/erthalion](https://github.com/erthalion)

 [@erthalion](https://twitter.com/erthalion)

 [dmitrii.dolgov at zalando dot de](mailto:dmitrii.dolgov@zalando.de)

 [9erthalion6 at gmail dot com](mailto:9erthalion6@gmail.com)