



NOSQL BEST PRACTICES

FOR POSTGRESQL



DMITRY DOLGOV

29-03-2018



Introduction

Less benchmarks

More opinionated best practices

Introduction

Application developers

DBAs

Extension developers

Application developers

When jsonb?

When jsonb?

→ You have a distinct flexible model

When jsonb?

- You have a distinct flexible model
- You need to work with data provided in document oriented format

When jsonb?

- You have a distinct flexible model
- You need to work with data provided in document oriented format
- Workaround for technical issues (large number of tables or expensive alignment)

When not jsonb?

When not jsonb?

→ Flexibility "just in case"

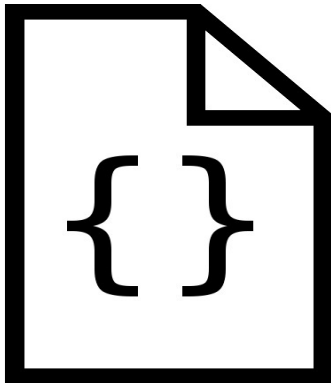
When not jsonb?

- Flexibility "just in case"
- Reluctance to create a migration

When not jsonb?

- Flexibility "just in case"
- Reluctance to create a migration
- Use jsonb column as a "garbage can"

Jsonb -> Relation



When to move from jsonb to relation?

When to move from jsonb to relation?

- Queries rely significantly in information about internal structure of documents

When to move from jsonb to relation?

- Queries rely significantly in information about internal structure of documents
- There are too many constraints for documents

When to move from jsonb to relation?

- Queries rely significantly in information about internal structure of documents
- There are too many constraints for documents
- Some parts of document are used much more frequently than other

```
SELECT id, created FROM some_table
WHERE
```

```
(data->>'name' = :a
AND (data @> ('{"items":[{"id":"' || :b || '"}]}'))
AND (data @> ('{"items":[{"elems":[{"name":"' || :c || '"}]}]}'))
AND (data @> ('{"items":[{"elems":[{"id":"' || :d || '"}]}]}'))
AND (data @> ('{"items":[{"name":"' || :e || '"}]}'))

ORDER BY created ASC, id ASC;
```

Complicated conditions

- jquery
- SQL/JSON

Complicated conditions

```
SELECT id, created FROM some_table  
WHERE
```

```
  data @@ 'items.#(id = '||:a||')'  
AND data @@ 'items.#.elems.#(name = '||:b||')'  
AND data @@ 'items.#.elems.#(id = '||:c||')'  
AND data @@ 'items.#(name = '||:d||')'
```

```
ORDER BY created ASC, id ASC;
```

Complicated conditions

```
SELECT id, created FROM some_table  
WHERE
```

```
  data @~ '$.items[*] ? (@id = '||:a||')'  
AND data @~ '$.items[*].elems[*] ? (@name = '||:b||')'  
AND data @~ '$.items[*].elems[*] ? (@id = '||:c||')'  
AND data @~ '$.items[*](@name = '||:d||')
```

```
ORDER BY created ASC, id ASC;
```

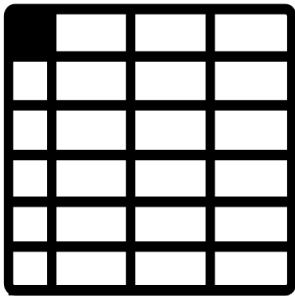
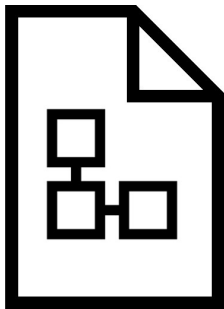
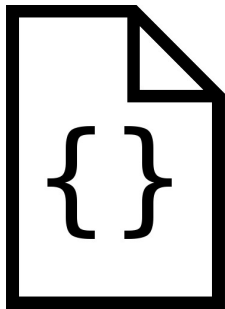
Complicated select

SELECT

```
st.data #>> '{item_a, another_item}' AS item_a,  
st.data #>> '{item_c}' AS item_c,  
jsonb_array_elements(  
    data #> '{item_b, subitem_a, subitem_b}'  
) ->> 'some_key' AS item_e
```

```
FROM some_table st LEFT JOIN another_table at  
ON (st.data #> '{item_b, key_a, key_b}') @>  
    jsonb_build_array(jsonb_build_object(  
        'key', 'some_key_name',  
        'value', at.data #>> '{item_b, another_item}'  
    ));
```

Jsonb -> Relation



Constraints

- Simple checks for value, type or size
- More convenient checks with jquery
- Json schema

Constraints

- Simple checks for value, type or size
- More convenient checks with jquery
- Json schema

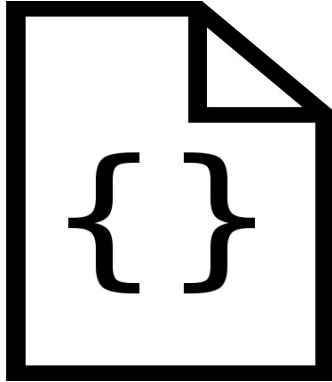
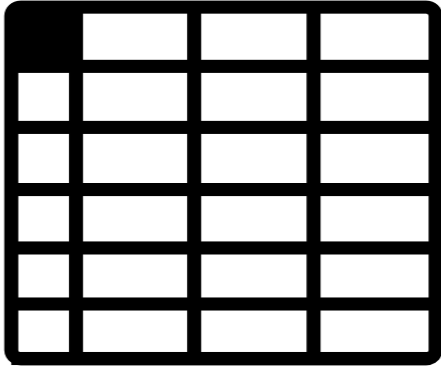
Constraints

```
CREATE TABLE test (  
  data jsonb,  
  CHECK (jsonb_typeof(data->'key') = 'array')  
);
```

```
CREATE TABLE test (  
  data jsonb,  
  CHECK (data @> 'key IS ARRAY OR key IS OBJECT')  
);
```

```
CREATE TABLE test (  
  data jsonb,  
  CHECK (validate_json_schema('{"key": "array"}', data))  
);
```

Relation -> Jsonb



```
SELECT jsonb_agg(query) FROM (  
    SELECT id, data  
    FROM jsonb_table  
) query;
```

Seamless interaction between json and rela- tion

```
[{
  "items": [
    {"id": 1, "value": "aaa"},
    {"id": 2, "value": "bbb"}
  ],
  {
    "items": [
      {"id": 3, "value": "aaa"},
      {"id": 4, "value": "bbb"}
    ]
  }
}]
```

```
WITH items AS (  
    SELECT jsonb_array_elements(data->'items')  
    AS item FROM test  
)  
SELECT * FROM items  
WHERE item->>'value' = 'aaa';  
  
item  
-----  
{ "id": 1, "value": "aaa" }  
{ "id": 3, "value": "aaa" }  
(2 rows)
```

```
WITH items AS (  
    SELECT jsonb_array_elements(data->'items')  
    AS item FROM test  
)  
SELECT * FROM items  
WHERE item->>'value' = 'aaa';  
  
item  
-----  
{ "id": 1, "value": "aaa" }  
{ "id": 3, "value": "aaa" }  
(2 rows)
```



```
{  
  "items": {  
    "item1": {"status": true},  
    "item2": {"status": true},  
    "item3": {"status": false}  
  }  
}
```

```
WITH items AS (  
    SELECT jsonb_each(data->'items')  
    AS item FROM test  
)  
SELECT (item).key FROM items  
WHERE (item).value->>'status' = 'true';
```

key

item1

item2

(2 rows)

```
WITH items AS (  
    SELECT jsonb_each(data->'items')  
    AS item FROM test  
)  
SELECT (item).key FROM items  
WHERE (item).value->>'status' = 'true';
```

key

item1

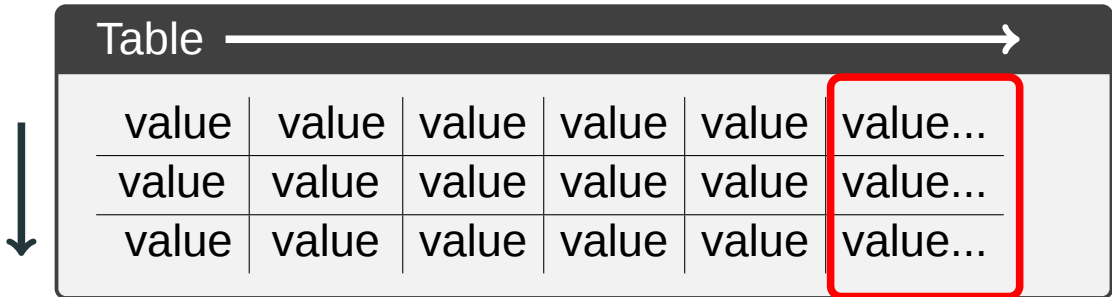
item2

(2 rows)

Multiple jsonb columns

- Keep at the end for readability
- tuple_deform (PG11, JIT compilation)

Multiple jsonb columns



The diagram illustrates a table structure with multiple jsonb columns. A dark grey header bar at the top is labeled "Table" on the left and features a white arrow pointing to the right. Below the header, a table is shown with three rows and six columns. Each cell in the table contains the text "value". The entire table, including the header bar, is enclosed in a light grey rounded rectangle. To the left of the table, a dark grey arrow points downwards. The sixth column of the table is highlighted with a red rectangular border.

Table →					
value	value	value	value	value	value...
value	value	value	value	value	value...
value	value	value	value	value	value...

Multiple jsonb columns

Table →

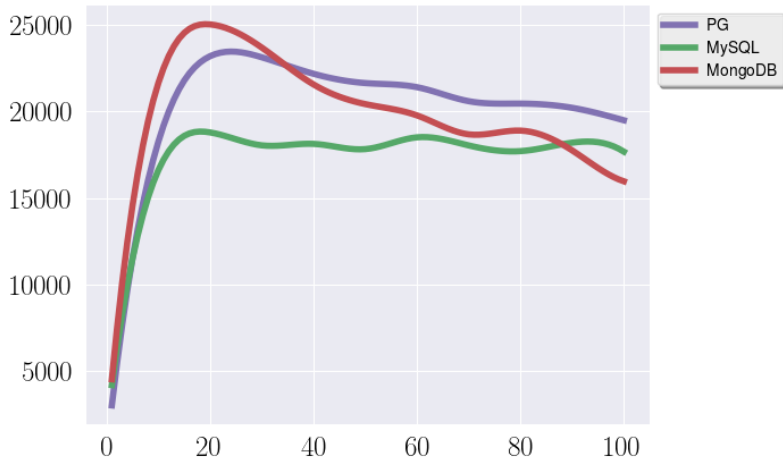
↓

value...	value	value	value	value	value
value...	value	value	value	value	value
value...	value	value	value	value	value

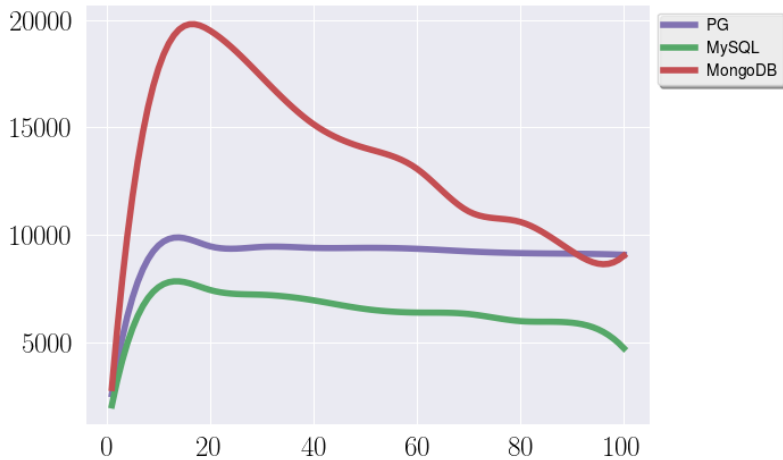
Document slice: in the DB or on the app?

- Amount of data passed from DB to application
- Performance hit in some cases
(multiple detoasting)

Throughput (ops/sec)



Throughput (ops/sec)



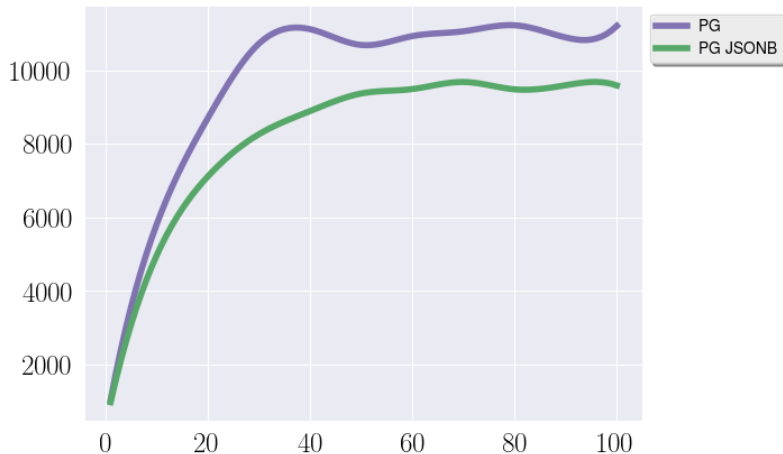
- Plain Json
- Binary Jsonb
- Relation

Insert workload

Throughput (ops/sec)

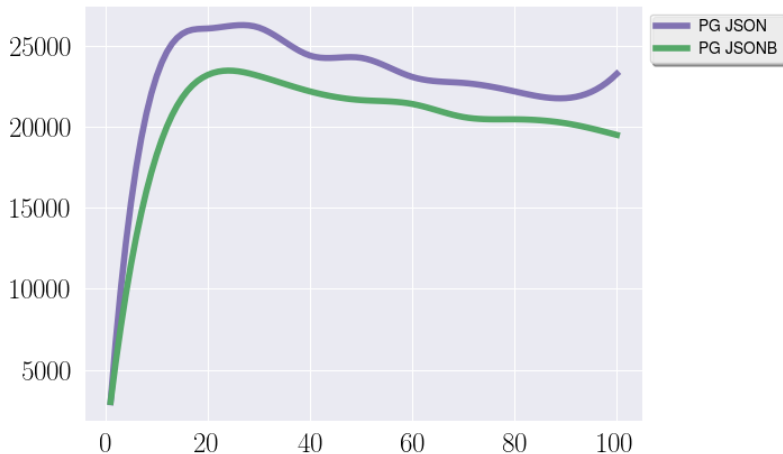


Throughput (ops/sec)

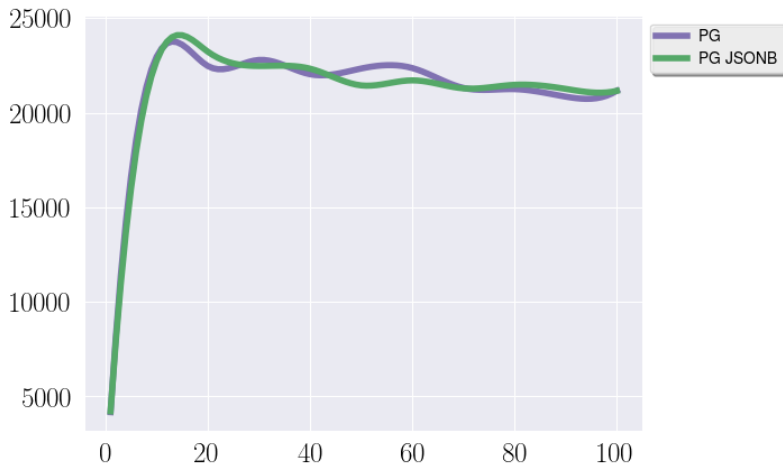


Read workload

Throughput (ops/sec)



Throughput (ops/sec)



Jsonb array vs regular array

- Store elements of different type?
Not really a "single model" idea.
- Syntax is less natural (this may change)
- Updates are slower
- Arrays are 1-based, Jsonb 0-based

```
SELECT array[0] FROM some_table;
```

```
SELECT jsonb->0 FROM some_table;
```

```
-- WIP
```

```
SELECT jsonb[0] FROM some_table;
```

```
UPDATE some_table SET array[0] = 'new_value';
```

```
UPDATE some_table
```

```
SET jsonb = jsonb_set(jsonb, '{0}', 'new_value');
```

```
-- WIP
```

```
UPDATE some_table SET jsonb[0] = 'new_value';
```

Some useful extensions

- jsquery
- postgres-json-schema
- is_jsonb_valid
- zson (jsonbc, custom compression methods WIP)
- jsonb_explorer

Types, please



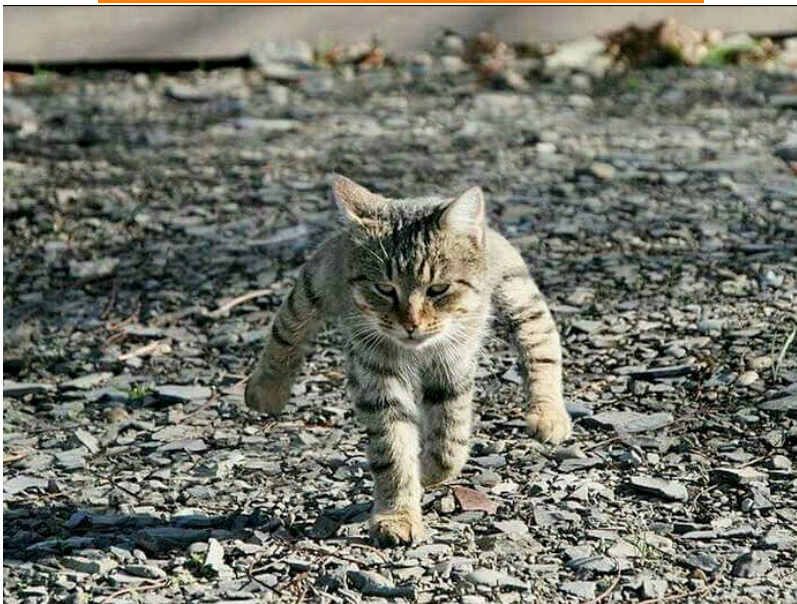
DBAs

Limitations

Size 256 MB

Depth - max_stack_depth

Stack depth is different for create & update



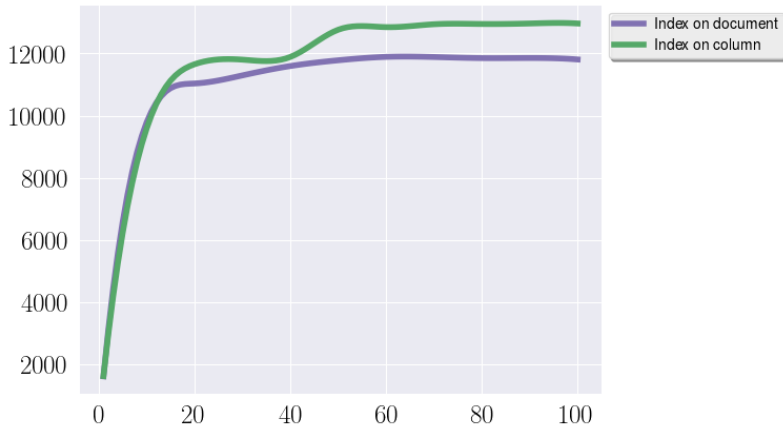
Indexing support

- GIN index (jsonb_ops, jsonb_path_ops)
- Functional BTree index
- jsquery strategies for GIN
- Partial indexes WIP

Place for ID

- Inside a document
- As a separate column

PG, Throughput (ops/sec)



Place for ID

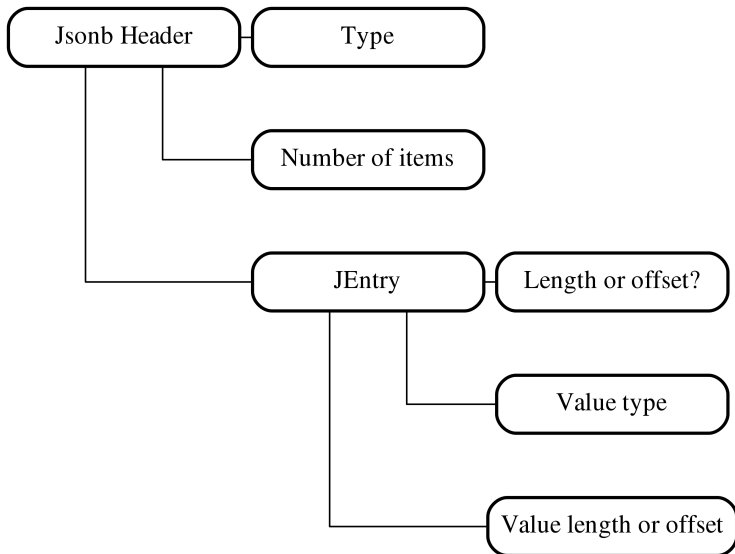
PostgreSQL 11 have HOT updates for some expression indexes, which will eliminate this problem.

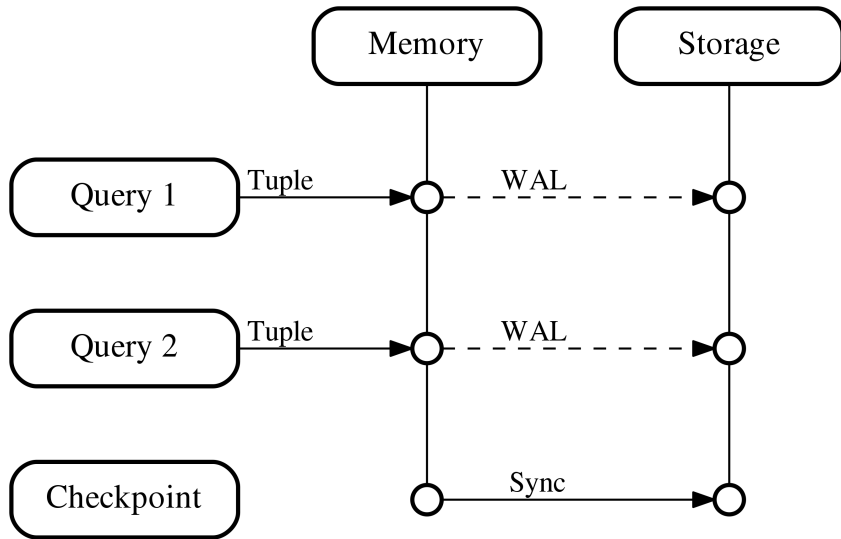
Statistics

- There is no proper selectivity estimation for jsonb
- Optimizer can give wrong estimations for GIN and complex queries
- Functional indexes

How much to write?

Jsonb vs Json

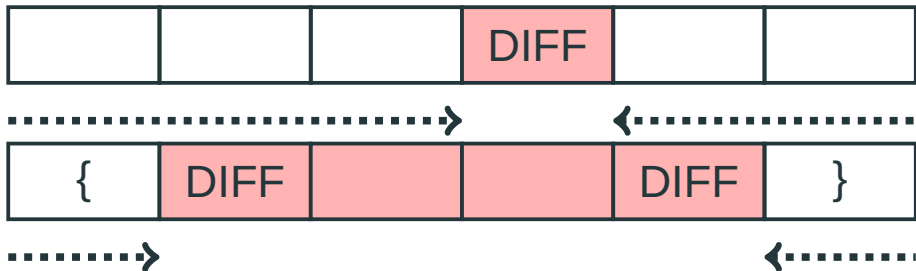




How much to write?

- Every update leads to update of an entire document (but it's ok)
- WAL can have a full document or just a diff
- Old and new tuples fit into the same page - diff
- Old and new tuples fit into the same page - full
- If logical decoding is enabled - full

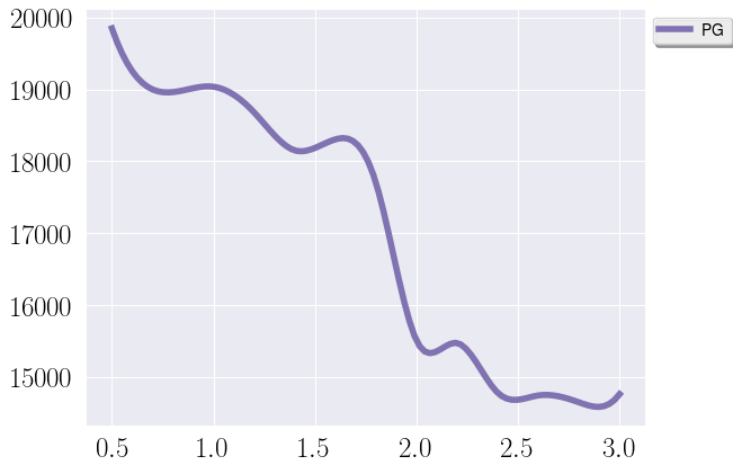
How much to write?



Huge documents

- TOAST has significant overhead (assemble, locks)
- Other than that linear degradation

Throughput, 40 clients



Alignment

Variable-length portion is aligned to a 4-byte

```
insert into test  
values('{"a": "aa", "b": 1}');
```

```
abaa\x20\x00\x00\x00\x00\x80\x01\x00
```

```
insert into test  
values('{"a": 1, "b": "aa"}');
```

```
\x00\x00ab\x00\x00\x20\x00\x00\x00\x00\x80\x01\x00aa
```

Extensions

Why to write an extension?

- Implement some convenient functionality (e.g. jsonb intersection)
- Create function optimized for your domain model

Why to write an extension?

`findJsonValueFromContainer` -> binary search
_id in a fixed position?

- Raw Jsonb container when search for an element
- Iterate through JsonbValue when update

Reuse infrastructure

- findJsonValueFromContainer
- JsonbIterator
- addToParseState
- worker functions


Random tips

- Clone iterator
- String are not null-terminated

Questions?

 github.com/erthalion

 [@erthalion](https://twitter.com/erthalion)

 9erthalion6 at gmail dot com